



Budapest University of Technology and Economics
Department of Artificial Intelligence and Systems Engineering

Artificial intelligence – VIMIAC16-EN, VIMIAC10

2024 Fall Semester

Dr. Gábor Hullám



Deep neural networks

- In the mid-2000s, there were hardly any articles on neural networks that were accepted
- 2004 – G. Hinton - CIFAR (Canadian Institute for Advanced Research)
 - New packaging for neural networks: deep learning
- 2006 - Hinton, Osindero, Yee-Whye Teh: A fast learning algorithm for deep belief nets
 - New breakthrough: teaching deep belief nets layer by layer,
 - Restarts research

Deep learning – ImageNet

2009 - J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li and L. Fei-Fei:
ImageNet: A Large-Scale Hierarchical Image Database

- high-resolution images, 14M images, 10k+ categories

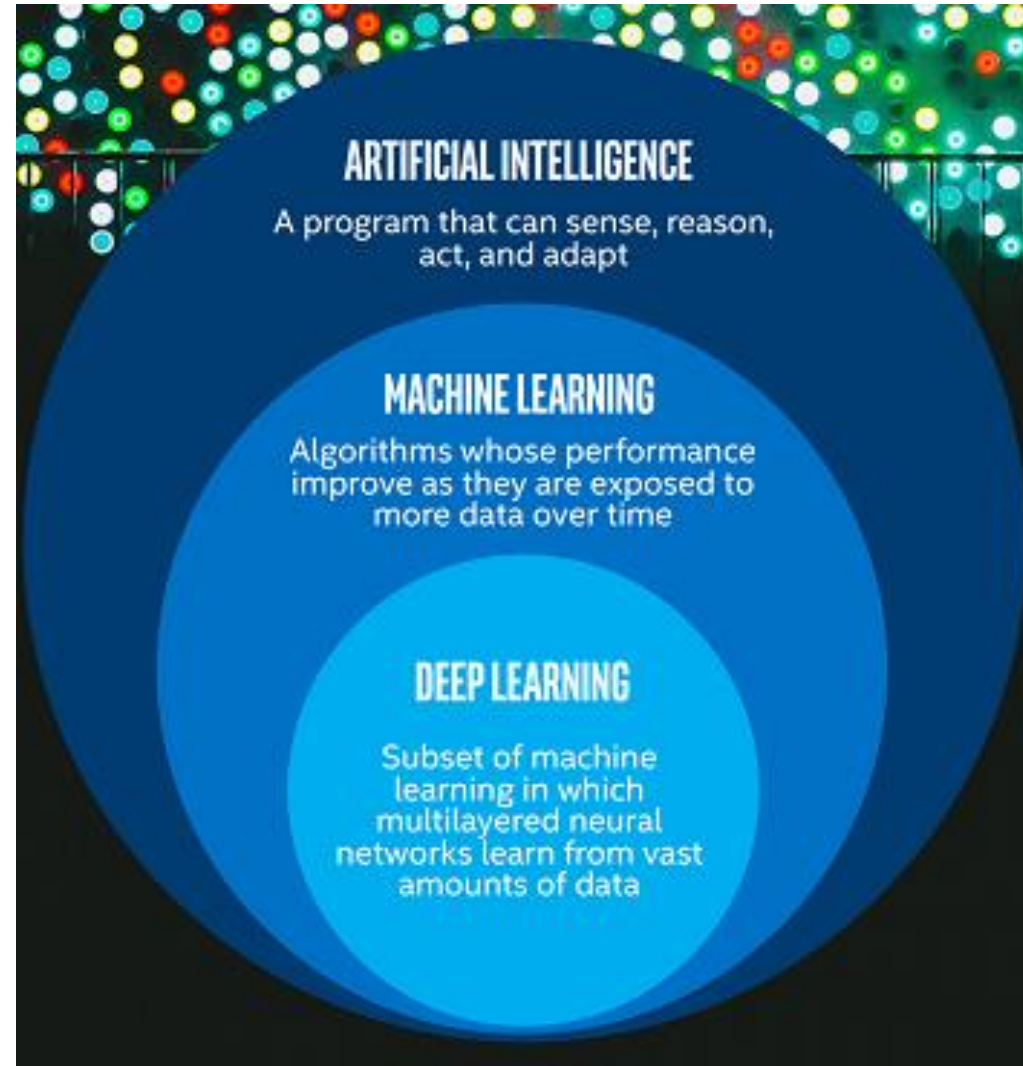
<http://www.image-net.org/about-stats>



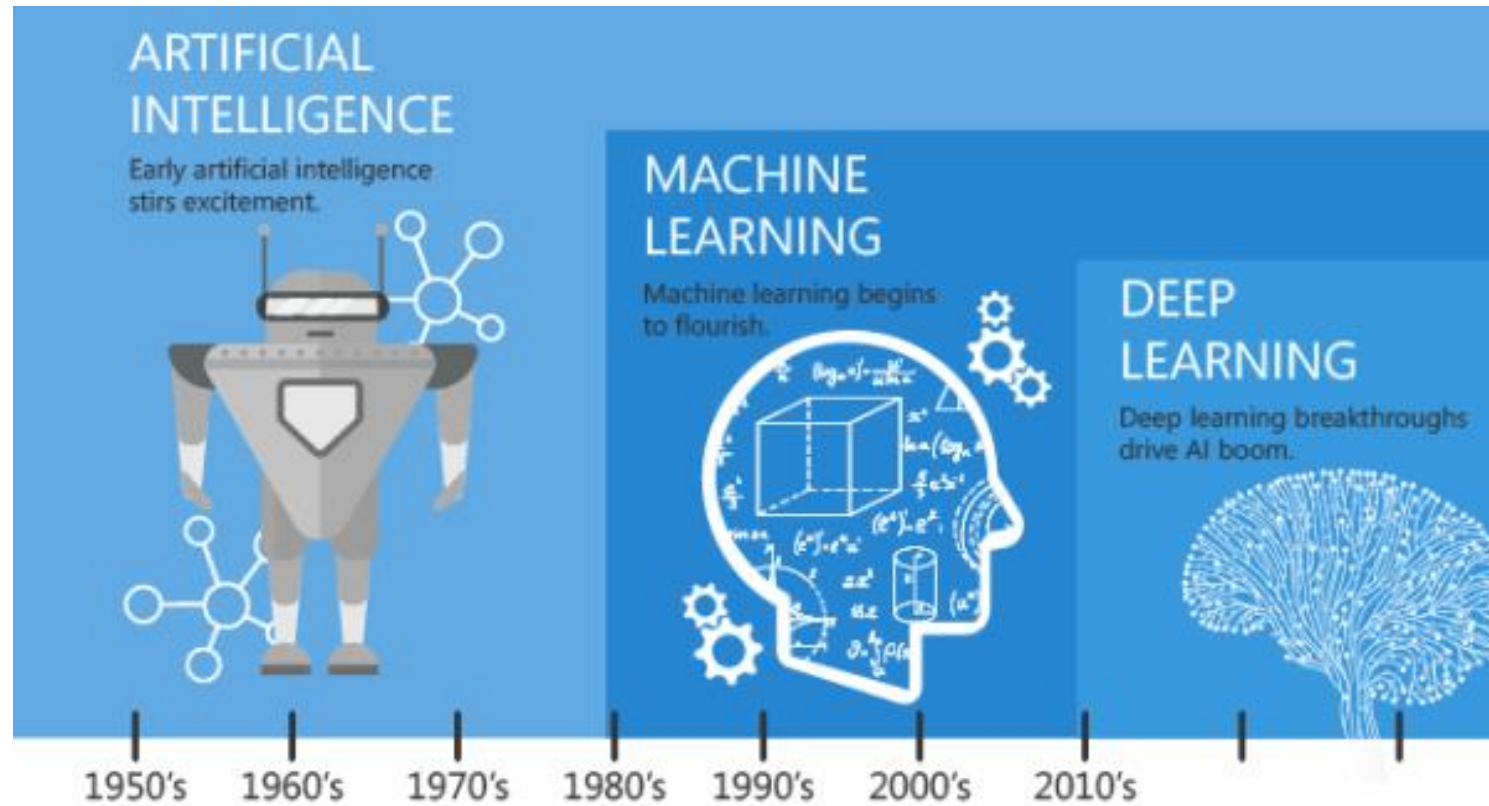
Deep learning - ILSVRC

- 2012 - Alex Krizhevsky, Ilya Sutskever, Geoffrey E Hinton: Imagenet classification with deep convolutional neural networks
 - This is the first purely neural network solution to achieve the best result of the year (SVM until then)
 - From there, deep learning becomes mainstream

AI – ML - DL




AI - ML - DL




<https://shawnennis.com/the-technology-of-machine-learning-with-ai>

AI – ML - DL



<https://www.cargroup.org/behind-headlines-artificial-intelligence-challenges-using-ai-automotive-industry/>

Background requirement: compute power




Moravec, H., 1998. When will computer hardware match the human brain. *Journal of evolution and technology*, 1(1), p.10.

Deep reinforcement learning

- 2013 – DeepMind: Playing Atari with Deep Reinforcement Learning
- 2016 – DeepMind: AlphaGo beats Go World Champion Lee Sedol

Go

- ▶ Google DeepMind
- ▶ Monte Carlo tree search
- ▶ 2016: 9 dan
- ▶ 2017: wins against human champion



ARTICLE

doi:10.1038/nature16961













Mastering the game of Go with deep neural networks and tree search

David Silver^{1*}, Aja Huang^{1*}, Chris J. Maddison¹, Arthur Guez¹, Laurent Sifre¹, George van den Driessche¹, Julian Schrittwieser¹, Ioannis Antonoglou¹, Veda Panneershelvam¹, Marc Lanctot¹, Sander Dieleman¹, Dominik Grewe¹, John Nham², Nal Kalchbrenner¹, Ilya Sutskever², Timothy Lillicrap¹, Madeleine Leach¹, Koray Kavukcuoglu¹, Thore Graepel¹ & Demis Hassabis¹

Deep learning - outcomes

■ Caption generation for images


- A. Karpathy, L. Fei-Fei: „Deep Visual-Semantic Alignments for Generating Image Descriptions”
- <https://research.googleblog.com/2014/11/a-picture-is-worth-thousand-coherent.html>

Describes without errors	Describes with minor errors	Somewhat related to the image	Unrelated to the image
 <p>A person riding a motorcycle on a dirt road.</p>	 <p>Two dogs play in the grass.</p>	 <p>A skateboarder does a trick on a ramp.</p>	 <p>A dog is jumping to catch a frisbee.</p>
 <p>A group of young people playing a game of frisbee.</p>	 <p>Two hockey players are fighting over the puck.</p>	 <p>A little girl in a pink hat is blowing bubbles.</p>	 <p>A refrigerator filled with lots of food and drinks.</p>
 <p>A herd of elephants walking across a dry grass field.</p>	 <p>A close up of a cat laying on a couch.</p>	 <p>A red motorcycle parked on the side of the road.</p>	 <p>A yellow school bus parked in a parking lot.</p>

Deep learning - outcomes

- Style 'learning'

- <https://github.com/jcjohnson/neural-style>
<https://imgur.com/gallery/4LTaQ>



Deep learning - outcomes

- Comprehension

- Xiong, W. et al. (2016). Achieving human parity in conversational speech recognition

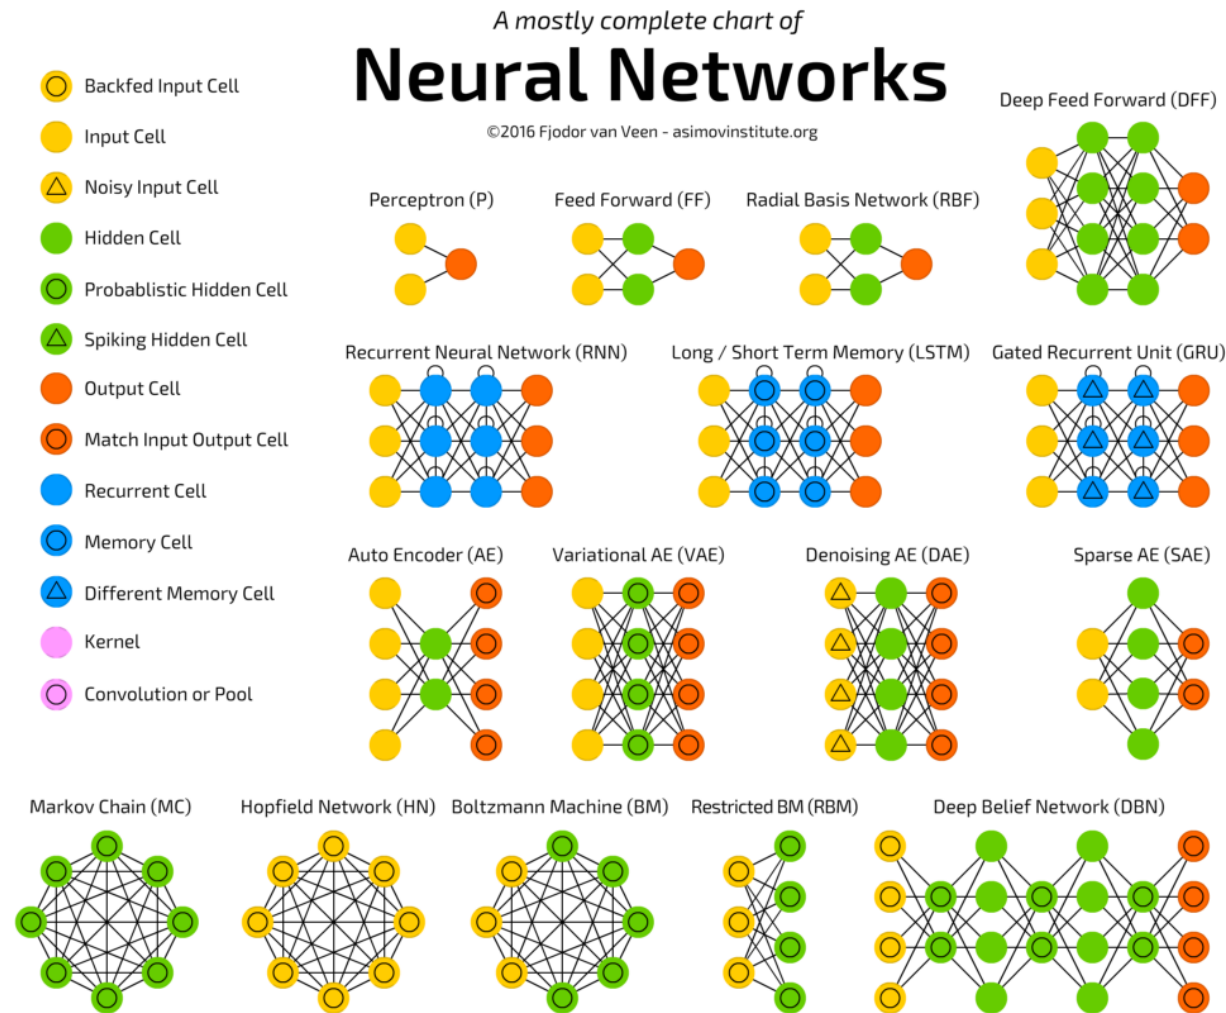
- Speech generation

- Oord, Aaron van den, et al. "Wavenet: A generative model for raw audio." arXiv preprint arXiv:1609.03499 (2016).

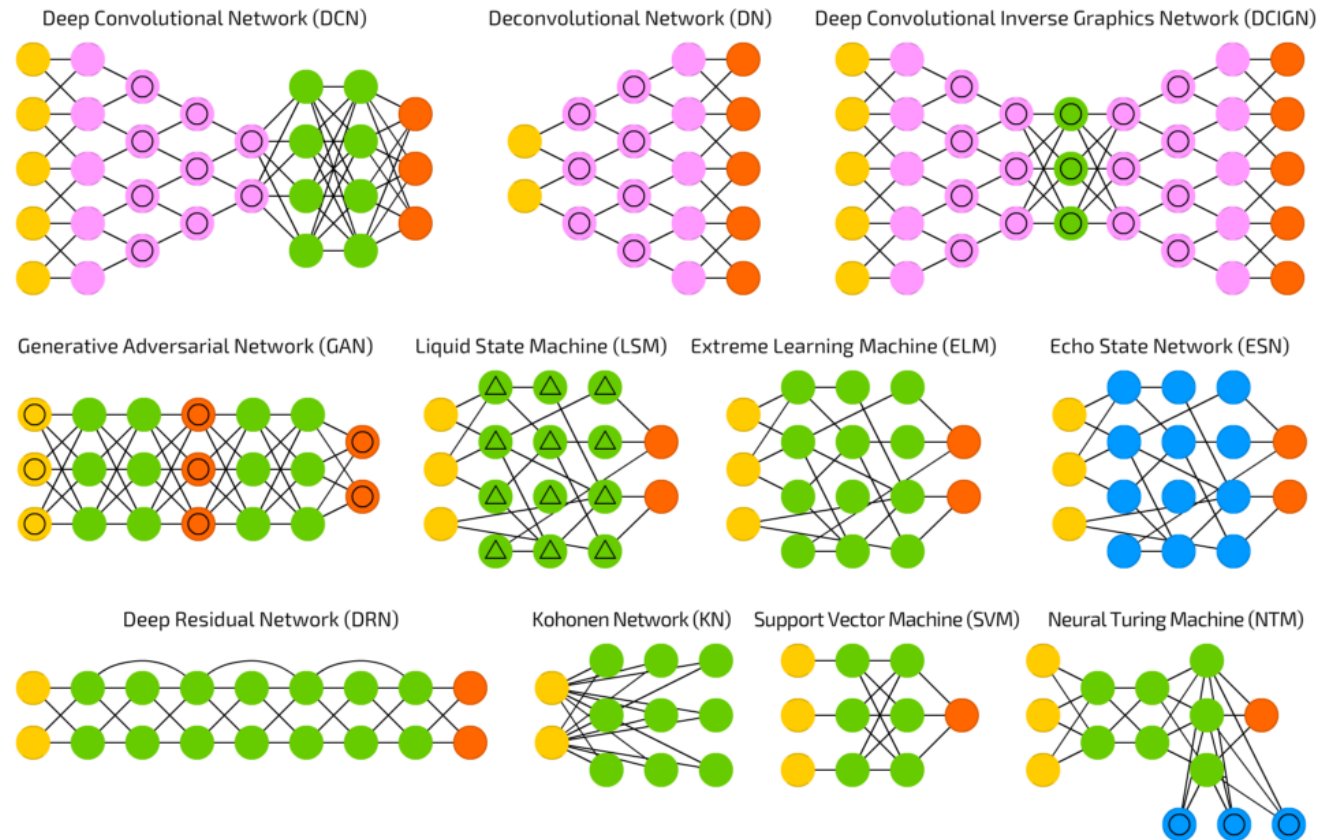
- Image classification

- C. Szegedy et al. (2014): GoogLeNet – „Going Deeper with Convolutions”

Types of neural networks




Types of neural networks 2.



What makes deep neural networks "better"?

- Humans also detect features of various levels of abstraction
- These from a **hierarchy** OR
- There is some form of **highlighting** = feature subset selection
- These principles should be applied
 - New action for compression / essence highlighting
 - Layer-by-layer weight training
 - Regularization

A Solution – Deep Belief Networks
(Hinton et al. 2006)




Alexnet 2012

ImageNet Classification with Deep Convolutional Neural Networks



2012



Convolution

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

⋮ ⋮

Each filter identifies a simple pattern(3 x 3).

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

Scalar
multiplication



1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

Convolution

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

3 -3

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

Convolution

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

-1	1	-1
-1	1	-1
-1	1	-1


Filter 2

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

For all filters



two 4 x 4 images from a
2 x 4 x 4 matrix

Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2


3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

-1	-1	-1	-1
-1	-1	-2	1
-1	-1	-2	1
-1	0	-4	3

CNN



cat dog



Convolution

Max Pooling

A new image

Convolution

Max Pooling

A new image

Flattened