


Performance Benchmarking



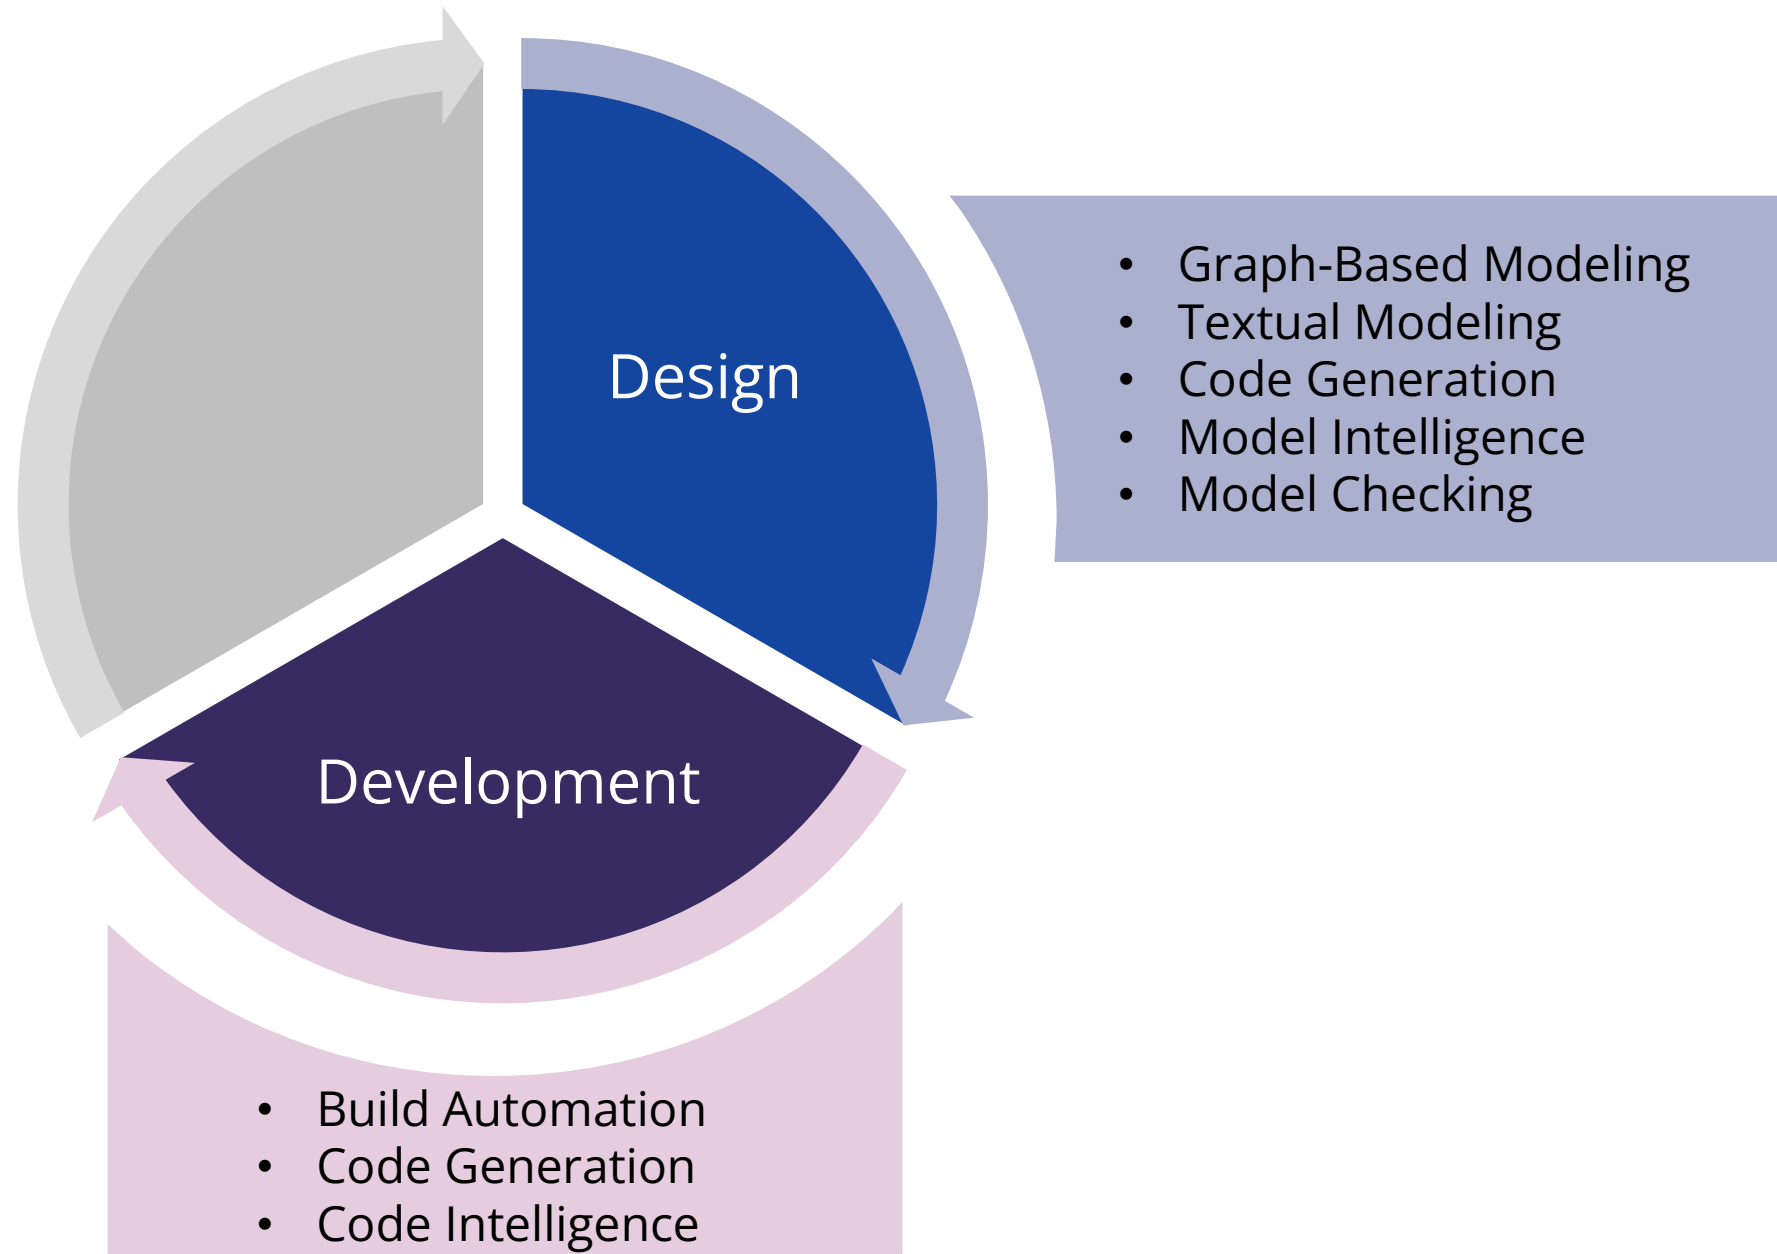
**Critical Systems
Research Group**

Overview




- Graph-Based Modeling
- Textual Modeling
- Code Generation
- Model Intelligence
- Model Checking

Overview



Overview

- Performance evaluation
- Data Analysis
- Code Quality
- Static Analysis
- Testing & Coverage



- Graph-Based Modeling
- Textual Modeling
- Code Generation
- Model Intelligence
- Model Checking

- Build Automation
- Code Generation
- Code Intelligence

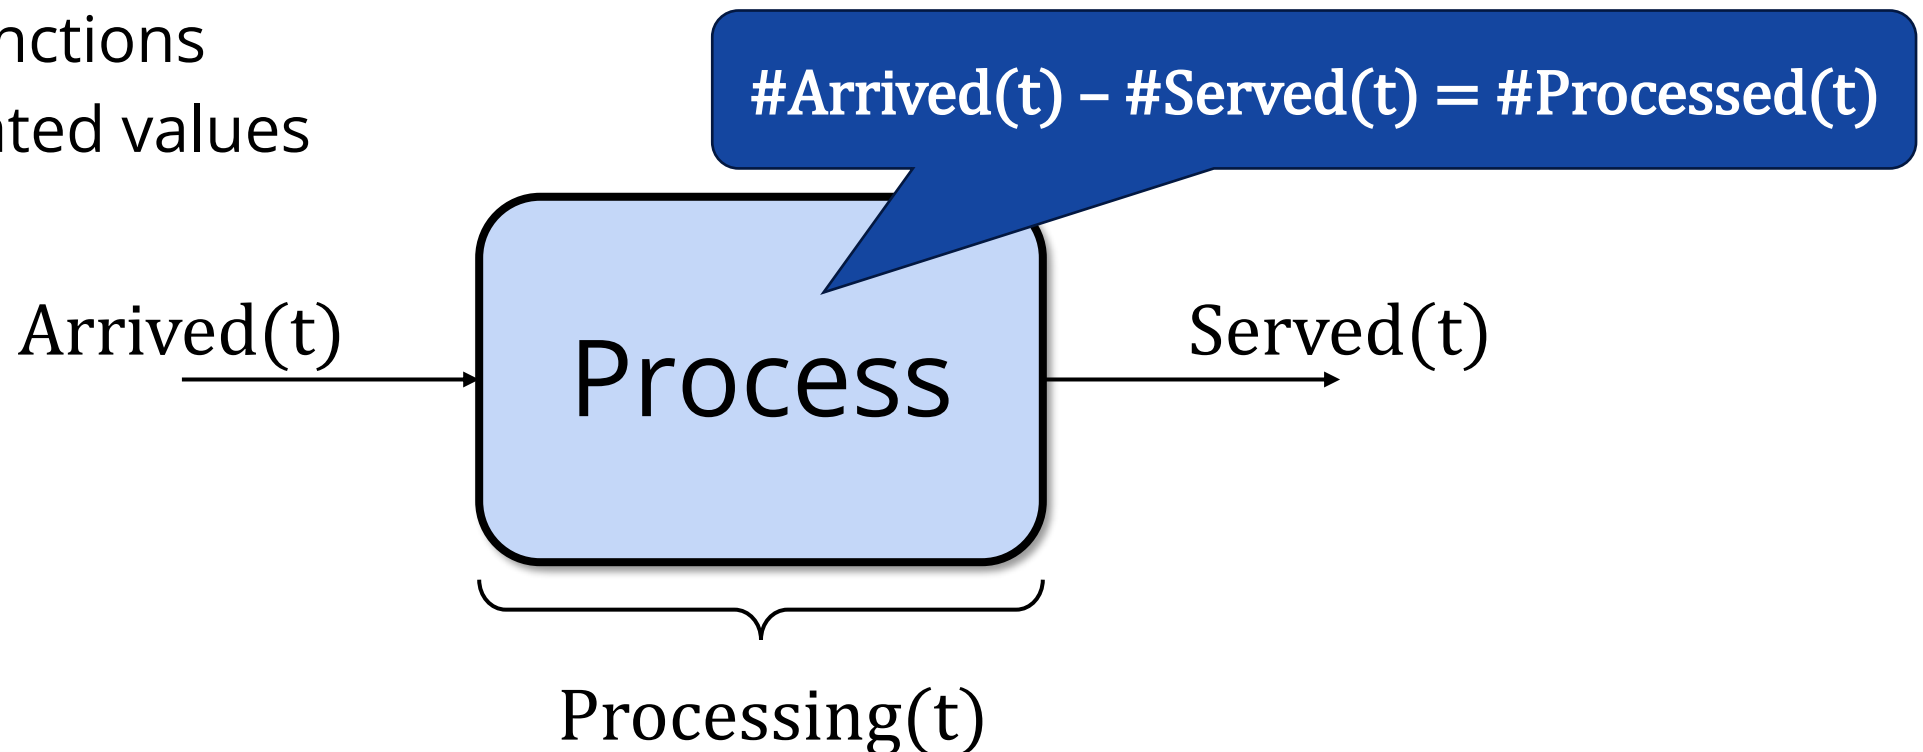
Performance Benchmark Overview

- Generic overview
 - Micro-benchmark
 - Macro-benchmark
-
- How to create a benchmark?
 - How to write an evaluation section?

Benchmarking Overview

Basic Model

- Executing a process multiple times
- Objectives: measure the time aspect of the behavior
- Describe the behavior with
 - Time functions
 - Aggregated values



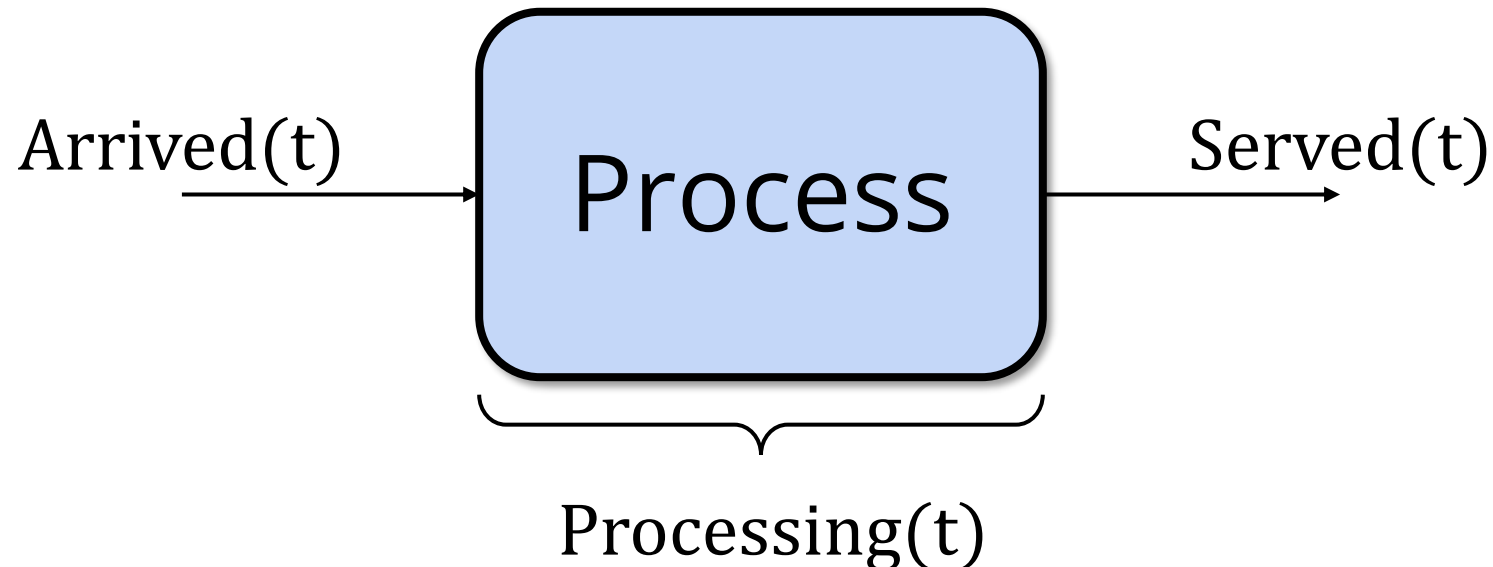
Definition: Arrival Rate, Throughput

- **Arrival rate:** number of **arriving** requests during a specific unit of time.

$$\lambda = \frac{Arrived(t)}{t} \quad [\lambda] = \frac{1}{s}$$

- **Throughput** (X_{PUT}): number of requests **processed** during a specific unit of time.




$$X = \frac{Served(t)}{t}$$



Definition: Stable State

- **Stable state:** *Under processing(t)* is approximately constant
 - Average values can be applied in such state!
 - A system is in balance, if:


$$\lambda = X$$



Definition: Stable State

- **Stable state:** *Under processing(t)* is approximately constant
 - Average values can be applied in such state!
 - A system is in balance, if:

In stable state:
Same number of logins
and logouts per minute



Limited Capacity – DoS

- N is not infinite in real life
- So, what is then?


Denial of Service Attack

Thursday, August 6, 2009 | By Biz Stone (@biz) 08/06/2009 - 15:00


Tweet

On this otherwise happy Thursday morning, Twitter is the target of a [denial of service attack](#). Attacks such as this are malicious efforts orchestrated to disrupt and make unavailable services such as online banks, credit card payment gateways, and in this case, Twitter for intended customers or users. We are defending against this attack now and will continue to update our [status blog](#) as we continue to defend and later investigate.


Load Diagram




Load Diagram




Load Diagram




Maximum Throughput




Utilisation




Approximative Load Function



Real Load Diagram



Real Load Diagram



Definitions

Maximum throughput: **maximal reachable** throughput

– Symbol: X^{\max}

Utilisation: **ratio** of the actual and the maximum throughput


– Symbol: $U = \frac{X}{X^{\max}}$

Thrashing: throughput decreases during overloaded operation


Effects of Load Fluctuation

- Average values vs. **real load**

In case of an overloaded system the waiting time can
2-3 magnitude higher



Outlook: Data Visualization



Micro- Benchmark

Micro-Benchmark Overview

- **Goal:** evaluate the performance of **small code fragments**
- Small but **important** piece of code

*We should forget about small efficiencies, say about 97% of the time:
premature optimization is the root of all evil"*

Donald Knuth

- In a highly controlled environment
 - Avoid optimization by compilers

Examples (from Refinery)

- We know that a hash function is critical in our application
 - Some hash functions are **slower** but more effective
 - Others are **quicker** but they clash more
 - How to evaluate the performance of each hash function?
- Which Map implementation should I use?
*Google Guava IntObjectMap **vs** good old java.util.HashMap?*
- New data structure, how to evaluate its performance?

How to execute those measurement precisely?

JMH: Java Microbenchmark Harness

- Library to execute precise Microbenchmarks
- It provides:
 - Measurement environment
 - Annotations to design measurements
 - Automated measurement execution
- It provides strict framework
- To avoid typical errors

Benchmark types

- We can annotate methods as benchmarks

```
@Benchmark
@BenchmarkMode(Mode.AverageTime)
public void method() { /* Do something */ }
```

- Supported modes:
 - **Throughput:** Operations per unit of time
 - **SampleTime:** Time distribution, percentile estimation
 - **SingleShotTime:** Time the single execution
- Example output:

Benchmark	Mode	Samples	Score-Idea	Score-Term	(MAX - MIN)/MAX	Units
measure	thrpt	25	341919594.645	344548681.629	0.007	ops/s

<https://github.com/artyushov/idea-jmh-plugin/blob/develop/research/results.md>

Warmup

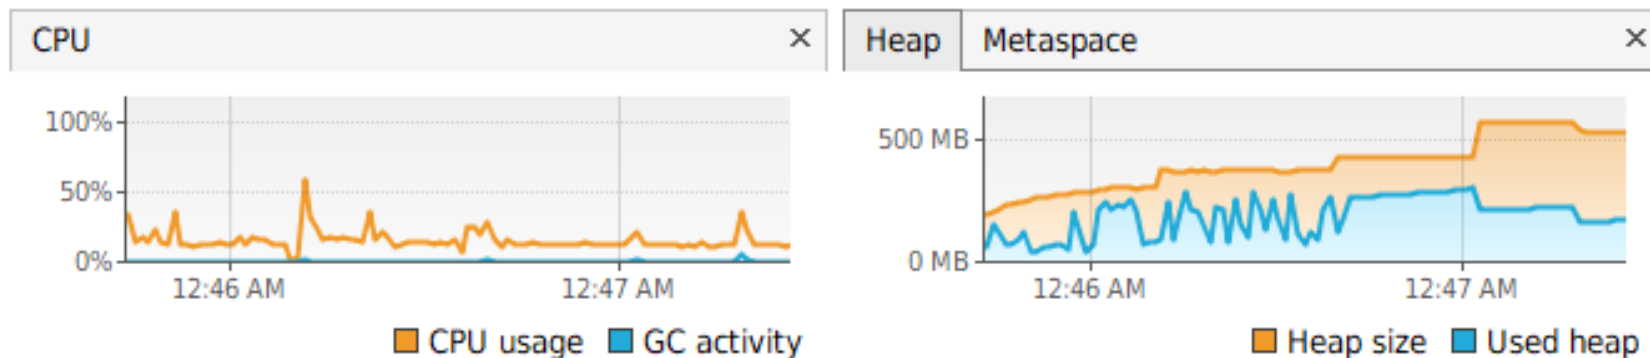
- In managed environment (e.g. Java), the runtime of a method **changes over time**
- We want to measure **peak performance**
- **Class loading:** for the **first time**, a lot of classes need to be loaded (e.g. Refinery is 102MB of pure *.class* files).
- **Code optimization:** managed environments (e.g., Java JIT) collect statistics during runtime and optimize code
→ better performance **after some iterations**

```
@Warmup(iterations = 20)
```





```
@Warmup(time = 1, timeUnit = TimeUnit.SECONDS)
```

Garbage collection

- In managed environments, GC is an automated process, which runs in the background.
- We need to avoid any interference.
- We need to measure the peak performance →
 - measure the system when there is enough memory, GC is not running
 - measure the system when in stable state, GC running in the background, evenly



How to measure memory

- How to measure memory requirements in a GC environment?
- Performance \Leftrightarrow Memory interaction.
- Lower the memory limit until the application crashes 😊.
 - **-Xmx16g** 
 - **-Xmx8g** 
 - **-Xmx4g** 
 - **-Xmx2g** 

Answer: memory requirement is between 2GB and 4GB

- This is a separate measurement.

Eliminating compiler optimization

- Calling the same function with the same parameters multiple times

```
public int calculateHash(Node node) { ... }
```

```
for (int i = 0; i < iterations; i++) {  
    calculateHash(Node node);  
}
```

- Compiler optimization, caching value →
After the second calculation, the runtime is ~0 ms.
- Solution: “black hole” object, that prevents caching

```
Blackhole blackhole = ...  
for (int i = 0; i < iterations; i++) {  
    blackhole.consume(calculateHash(Node node));  
}
```

Comparing multiple variants

- Run multiple measurements with multiple variations
- Like parametric tests

```
@State(Scope.Benchmark)  
public class ImmutablePutExecutionPlan {  
  
    @Param({ "100", "10000" })  
    public int nPut;  
  
    @Param({ "32", "1000", "100000" })  
    public int nKeys;  
  
    @Param({ "2", "3" })  
    public int nValues;  
    ...  
}
```


Macro- Benchmark

Macro-Benchmark Overview

- **Goal:** evaluate the performance of **an algorithm**
- The **whole application** is evaluated.
- **Sub-Goal 1:** discover and identify bottlenecks


The same relevant quote:

*We should forget about small efficiencies, say about 97% of the time:
premature optimization is the root of all evil”
Donald Knuth*

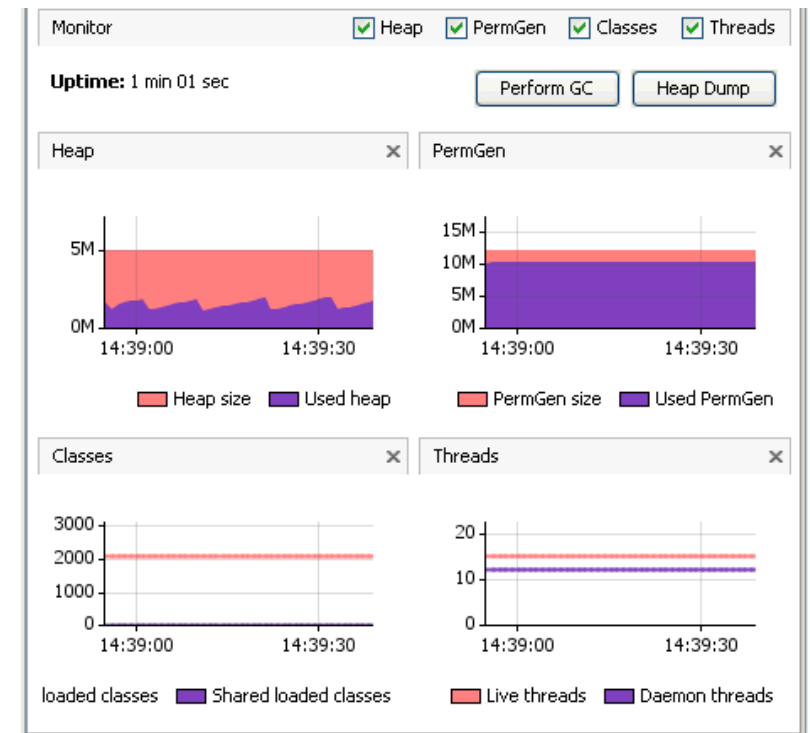
Sub-Goal 2: check the performance in increasing problem size

Profiling: identifying bottlenecks

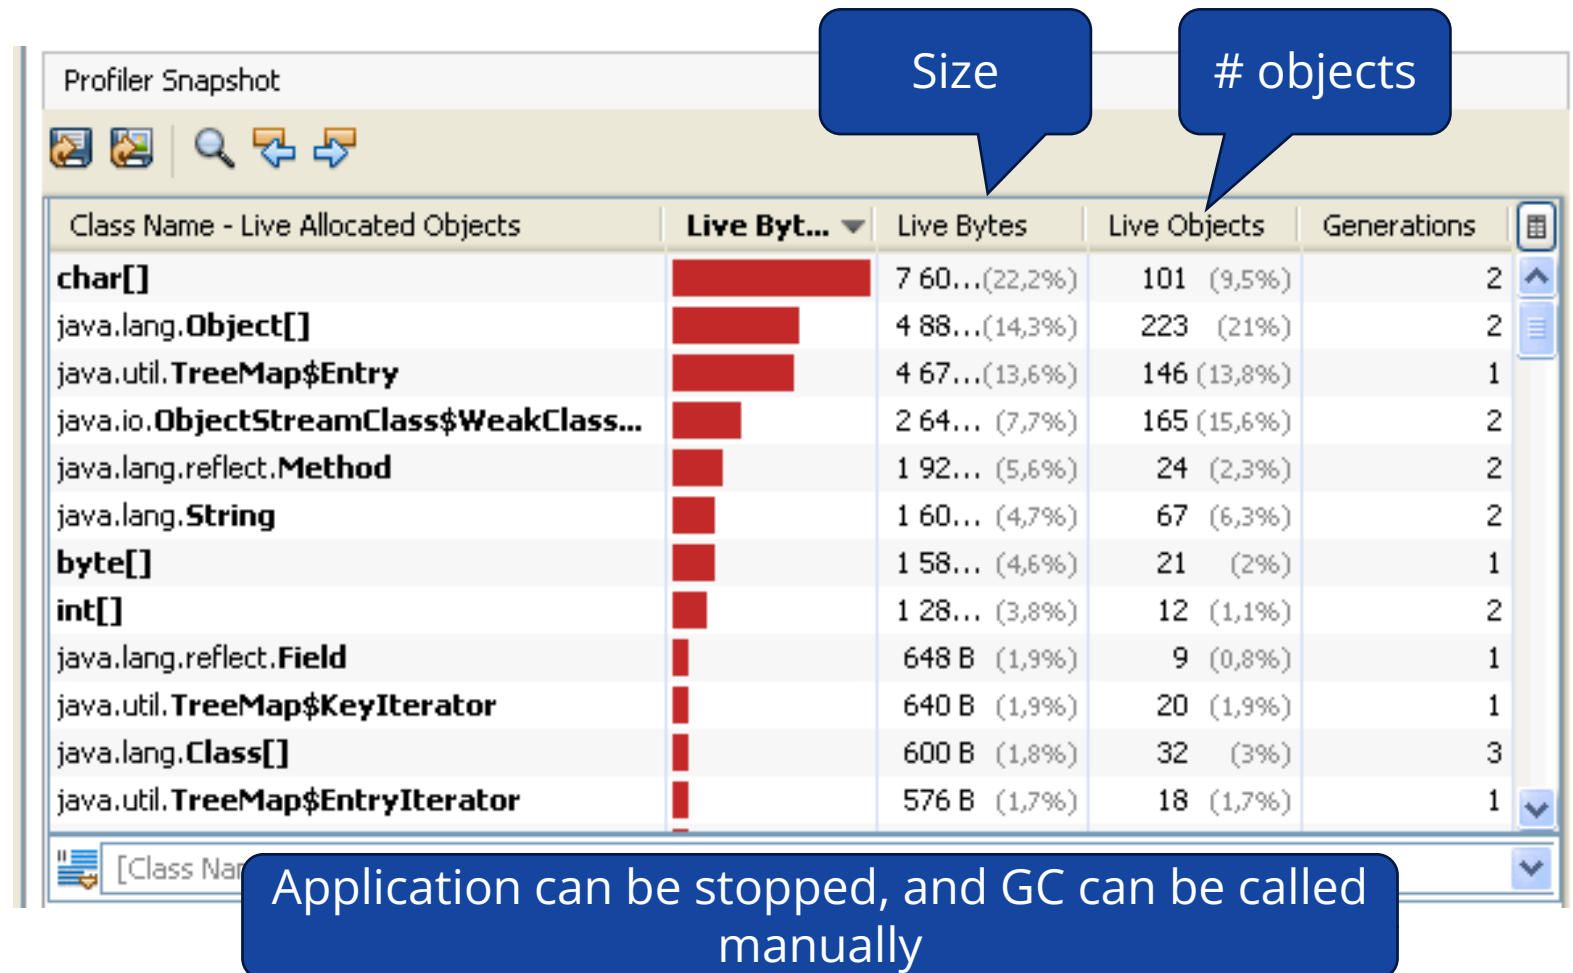
- VisualVM is a lightweight profiling tool
- Attaches itself to an application, and measures



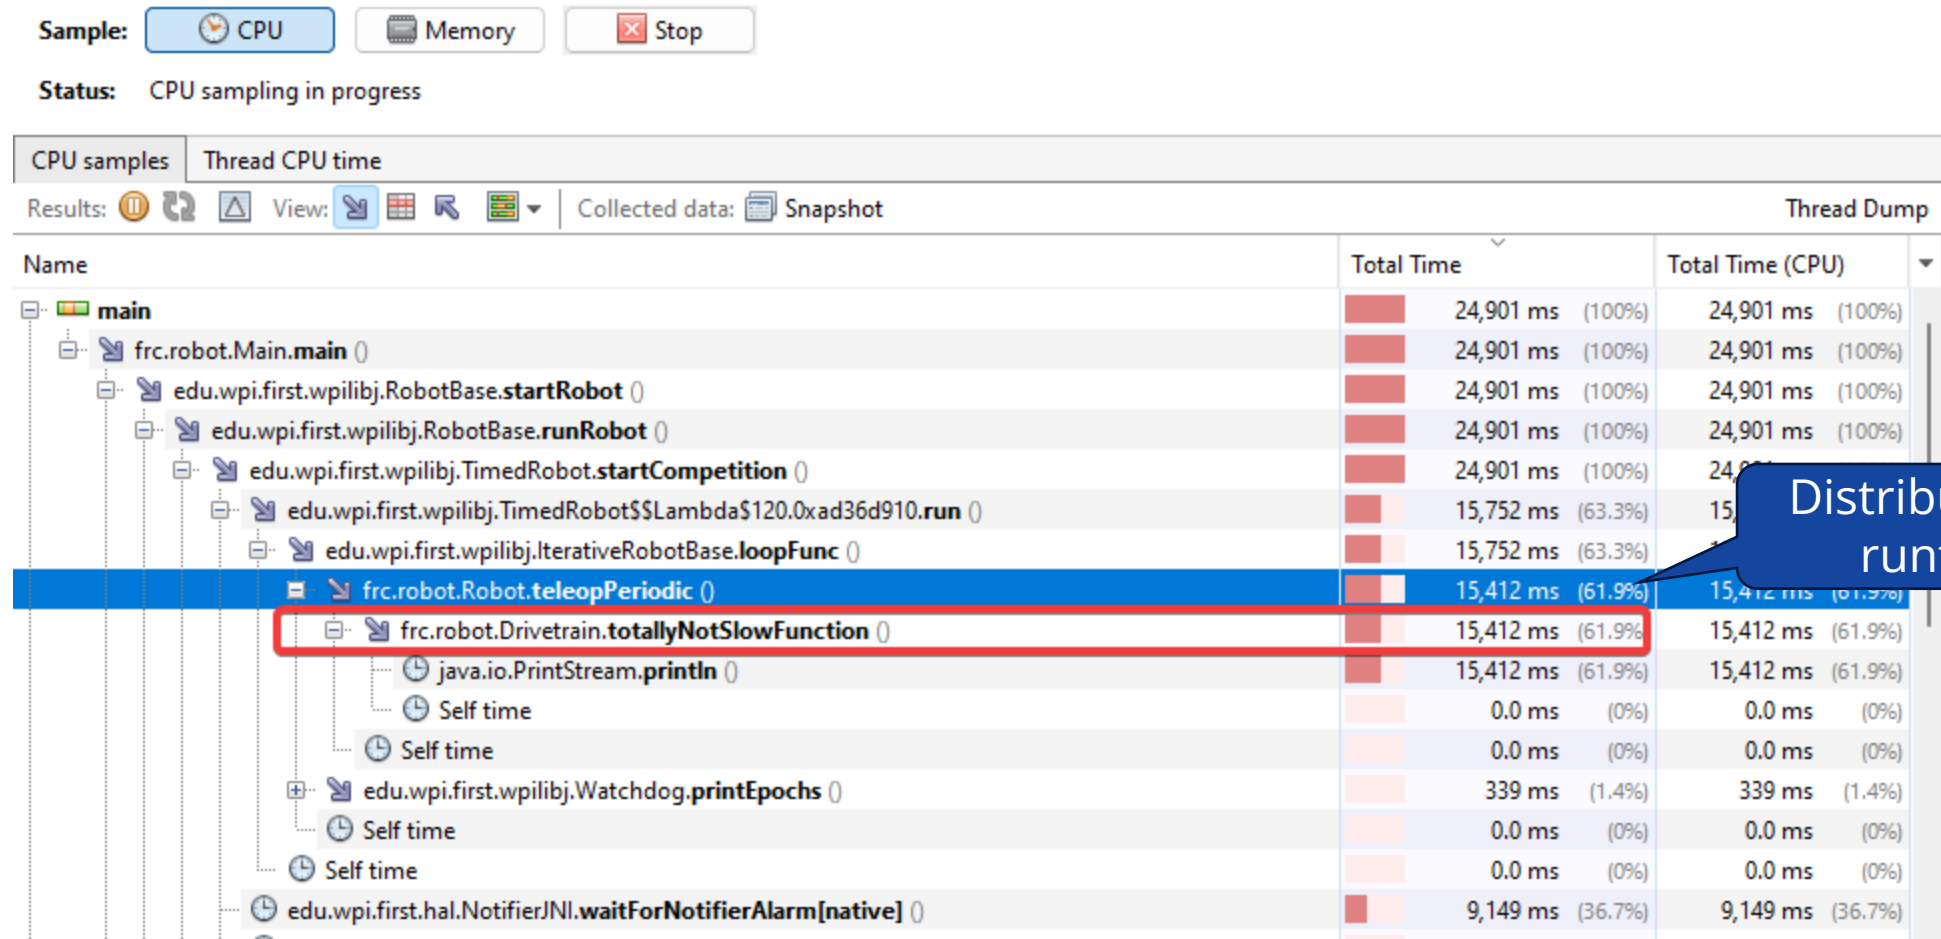
- Saves and visualizes measurements



Memory profiling



Performance profiling



<https://docs.wpilib.org/en/stable/docs/software/advanced-gradlerio/profiling-with-visualvm.html>

Two supported measurement method

- **Sampling: periodically** checks the stack and **makes snapshots**
 - Statistics are calculated on the snapshots
 - Hides fast methods, statistical guarantees
 - But realistic distribution
 - Typically the good solution
- **Profiling: changes** the code and **reports each call**.
 - Precisely captures the runtimes
 - But changes the code → we are not measuring the target application
 - Huge performance decrease, not realistic runtime.


How to create a Benchmark?

Important aspects of benchmark

- **Relevant:** realistic data
- **Independent:** multiple implementations can use it
- **Scalable:** small problem \Leftrightarrow big problem
- **Portable & Reproducible:** the measurement can be executed

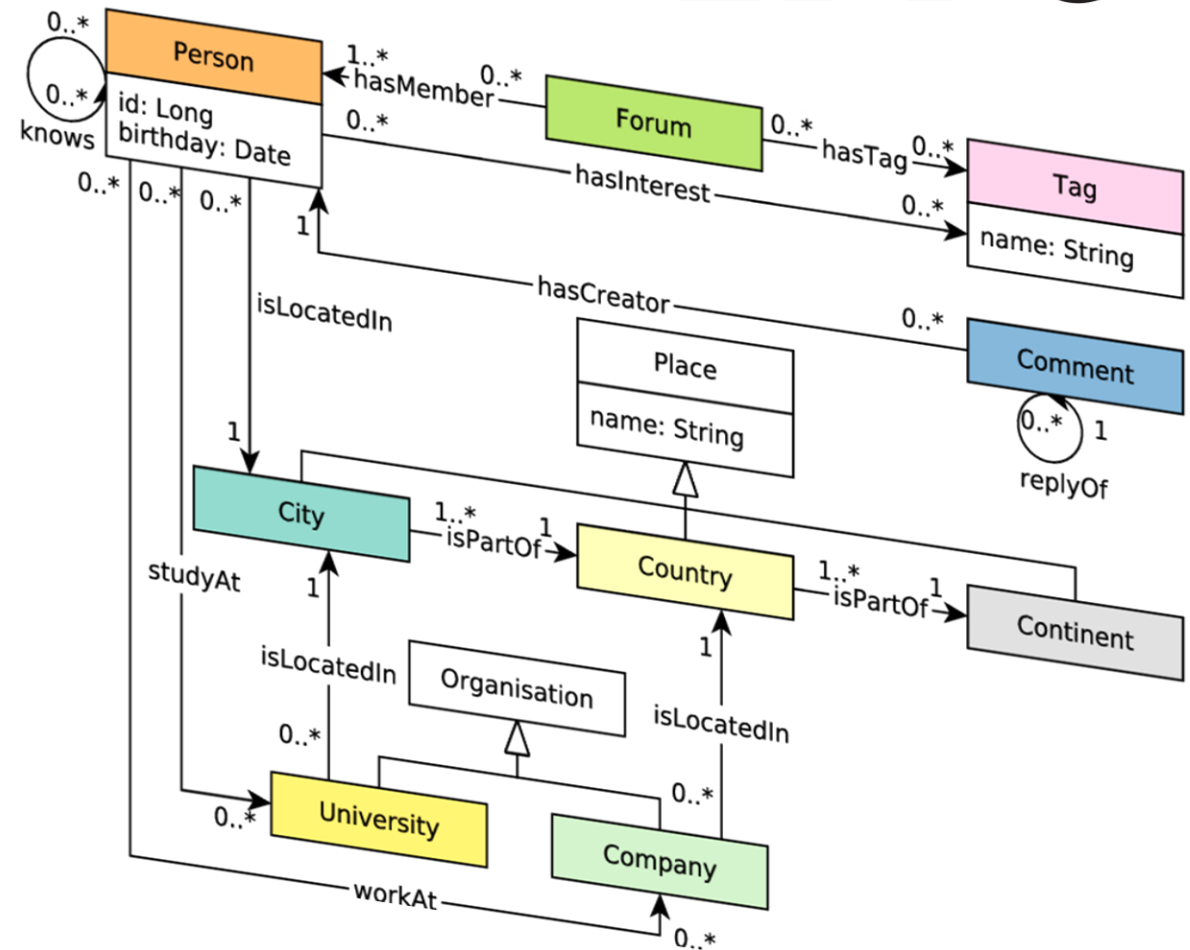
Example: BenchExec

- SV-COMP 2024: model checking benchmark
- Automated evaluation on large set of problems, competitors



LDBC benchmark

- Graph query benchmark
- Bigger and bigger graphs vs
- Complex queries
- Audited benchmark (\$)



How to Document?

Typical documentation of evaluation

- This can be a research report, TDK, thesis work, research paper
- There are typical sections

1. Research Questions.

Simple, straightforward question. Not yes-no.

- RQ1** How does our graph solver scale (in time and model size) when generating consistent models of increasing size?
- RQ2** How does our approach scale (in time and model size) compared to the widely used model finder Alloy [21]?
- RQ3** How do the different steps of the exploration influence performance of the graph solver?

Measurement environment

2. Selected domains

Selected example domains should be independent and representative.

3. Benchmarking setup

How is the evaluation executed?

Machine specification, software configuration, steps.

Communicating the results

4. Experimental results

Figures are explained here. Just simple explanation.



5. Research results

RQ2. According to the results (see Figure 8e–8j and Figure 8k), our approach scales much better as it generates models 1-2 orders of magnitude larger than Alloy could handle regardless of the back-end SAT solver which only had little impact on scalability.

Threats to validity

6. Threats to validity

- **Internal:** confidence in measurements
- **External:** confidence in generalization
- **Construct:** confidence in the relevance of measured metrics

Summary

Summary

- Generic overview
 - Micro-benchmark
 - Macro-benchmark
-
- How to create a benchmark?
 - How to write an evaluation section?