

Data-driven systems

Design patterns in the data-access layer
Concurrency handling
Object-relational mapping

Role of the data access layer

- Provide a high-level abstraction for data access
- Basic operations for data manipulation
- Concurrency handling

Object-relational mapping

Problem statement

Modelling

- Business Layer
 - > Object oriented modelling
 - > UML
 - > Design Patterns
 - > Not only data, but processes too
- Data Layer
 - > Entity-Relationship diagram
 - > UML data modelling profile
 - > Static, attribute-based

Purpose of ORM

- Object Relational Mapping
 - > Mapping business entities to the relational data model
 - > Connect data storage and business workflows
- Problems
 - > Different concepts
 - > Inheritance
 - > Shadow information
 - > Relations

Sources used

- Examples are from:
 - > <http://www.agiledata.org/essays/mappingObjects.html>


Object-relational mapping

Basic concept


Approach

- Basic idea
 - > Class → Table
 - > Field → Column
 - > Relationships → FK

Object-oriented






Relational schema




Problems

- Compound fields
 - > Customer
 - Address (Zip, City, Street)
- Different data types
 - > Conversion!




string

nvarchar(???)



Shadow information

- Required for persistency
 - > Keys (primary keys)
 - > Timestamps (for optimistic concurrency handling)
- Has no place in the business entity, but need to put it somewhere

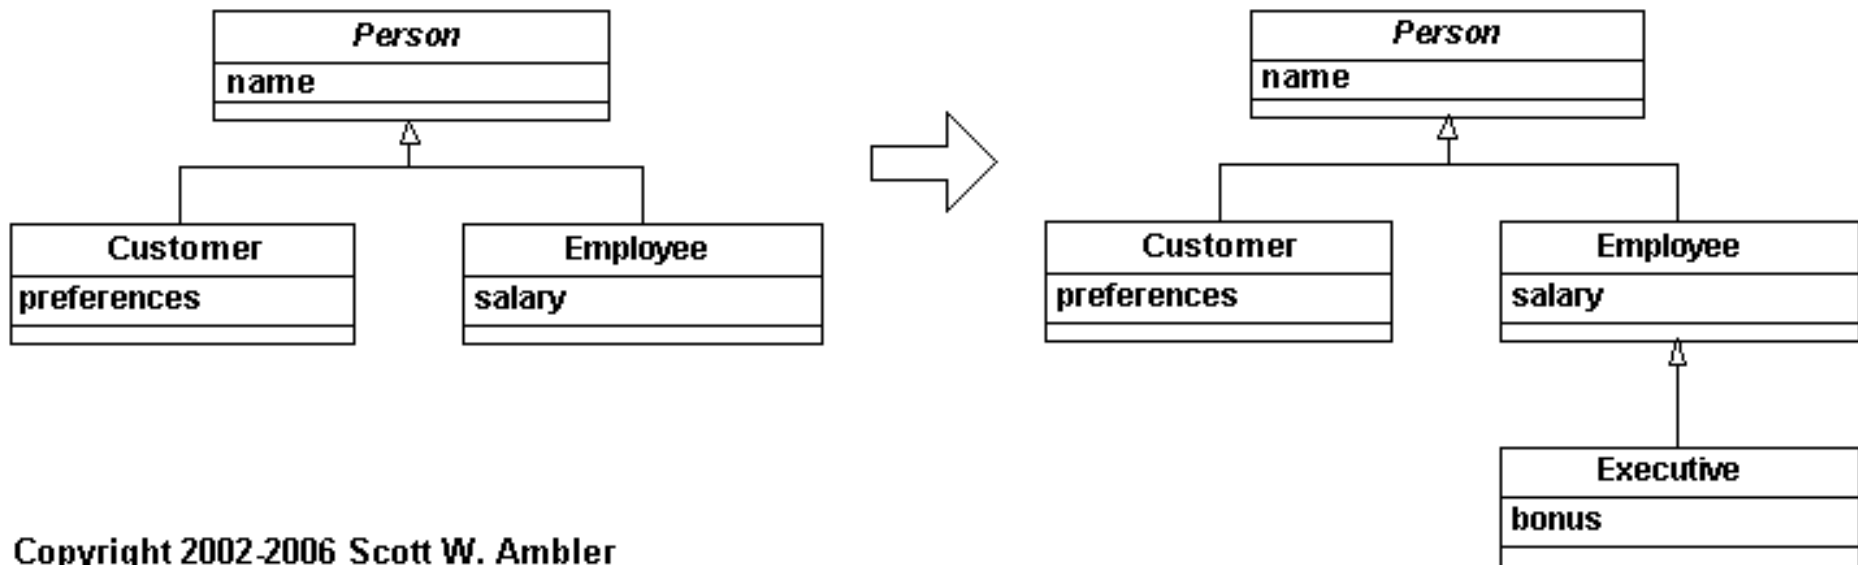


Object-relational mapping

Inheritance

Inheritance – Example

- Abstract class *Person*
- Multiple implementation
- New function built into the application -> new inheritance: *Executive*



Copyright 2002-2006 Scott W. Ambler

Inheritance

- Map the entire hierarchy into one table
- Map non-abstract types into their own table
 - > From which objects can be instantiated
- Map all types into their own table
 - > Abstract classes too
- General mapping independent of actual model

Mapping into a single table - 1

- List all attributes in the whole hierarchy
- Identify the type of the instance
 - > Encoded in a new column
 - > IsCustomer, IsEmployee,... columns
- Handling modification
 - > Add new columns

Mapping into a single table - 2

Person table

ID	Name	Prefer.	Salary	IsCustomer	IsEmployee
1	X Y	NULL	1234	0	1
2	W Z	xxxxx	NULL	1	0



Introduce Executive

ID	Name	Prefer.	Salary	IsCustomer	IsEmployee	Bonus	IsExecutive
1	X Y	NULL	1234	0	1	NULL	0
2	W Z	xxxxx	NULL	1	0	NULL	0
3	Q Q	NULL	456	0	1	999	1

Mapping into a single table - 3

- Advantages
 - > Simple
 - > Easy to add a new inheritor
 - > Easy to change the role of an instance
 - Employee → Executive
 - Employee and Customer too
- Drawbacks
 - > Wasting storage space
 - > Change in a single class changes storage model of all
 - > Hard to overview in case of a complex structure
- Use when ... - *design time* dimension
 - > Simple hierarchy

Mapping real classes - 1

- Unique table per class
- All attributes of the class
- Instance identifier
- Handling change
 - > New class → new table
 - > Change in an attribute → Must be propagated in the hierarchy

Mapping real classes - 2

Customer table

ID	Name	Prefer.
2	W Z	xxxxx

Employee table

ID	Name	Salary
1	X Y	1234



Introduce Executive

Executive table

ID	Name	Salary	Bonus
4	QQ	456	999

Mapping real classes - 3

- Advantages
 - > Intuitive
 - > Better fit to the OO paradigm
 - > Fast data access
- Drawbacks
 - > In case a class is modified → Might affect multiple tables
 - > Instances that have multiple roles
 - Employee → Executive
 - Employee and Customer at the same time
- Use when ... - *design time* dimension
 - > Structure that changes infrequently

Mapping all classes - 1

- Tables match the inheritance hierarchy
- Map parent-child relationship with foreign keys
- Instance identifier

Mapping all classes - 2

Person table

ID	Name
1	XY
2	WZ

Customer table

CID	PID (FK)	Prefer.
11	2	xxxxx

Employee table

EID	PID (FK)	Salary
21	1	1234



Introduce Executive

Executive table

XID	EID (FK)	Bonus
41	22	999

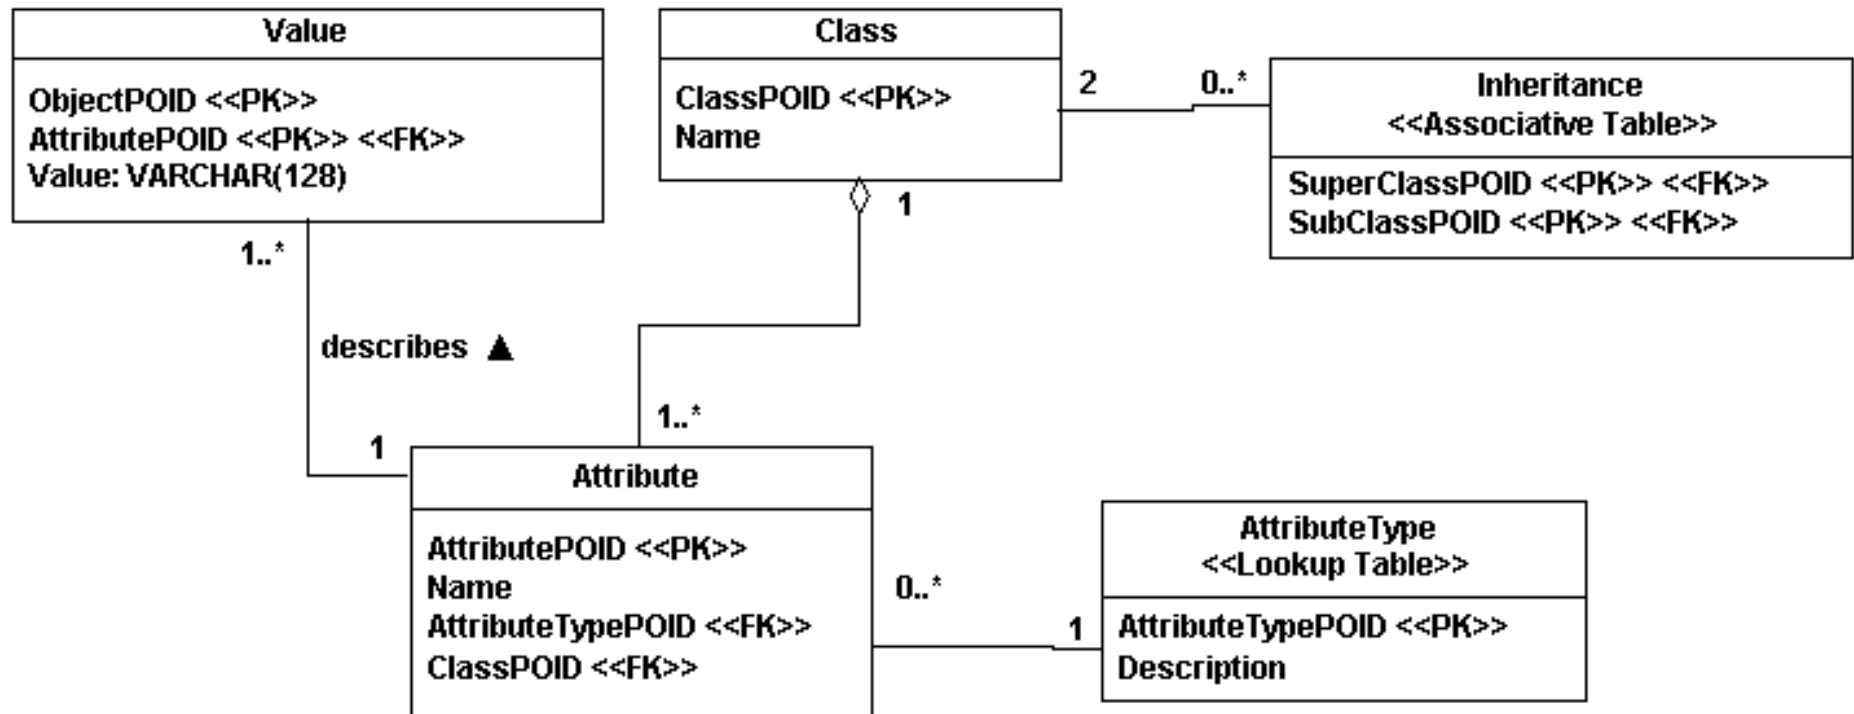
Mapping all classes - 3

- Advantages
 - > Intuitive
 - > Easy to modify parent class structure
- Drawbacks
 - > Complex database scheme
 - > Data of a single instance is scattered in multiple tables
 - Requires more complex queries
 - Join is needed → slower
- Use when ... - *design time* dimension
 - > Complex hierarchies
 - > Structure that changes frequently

Mapping into a general scheme - 1

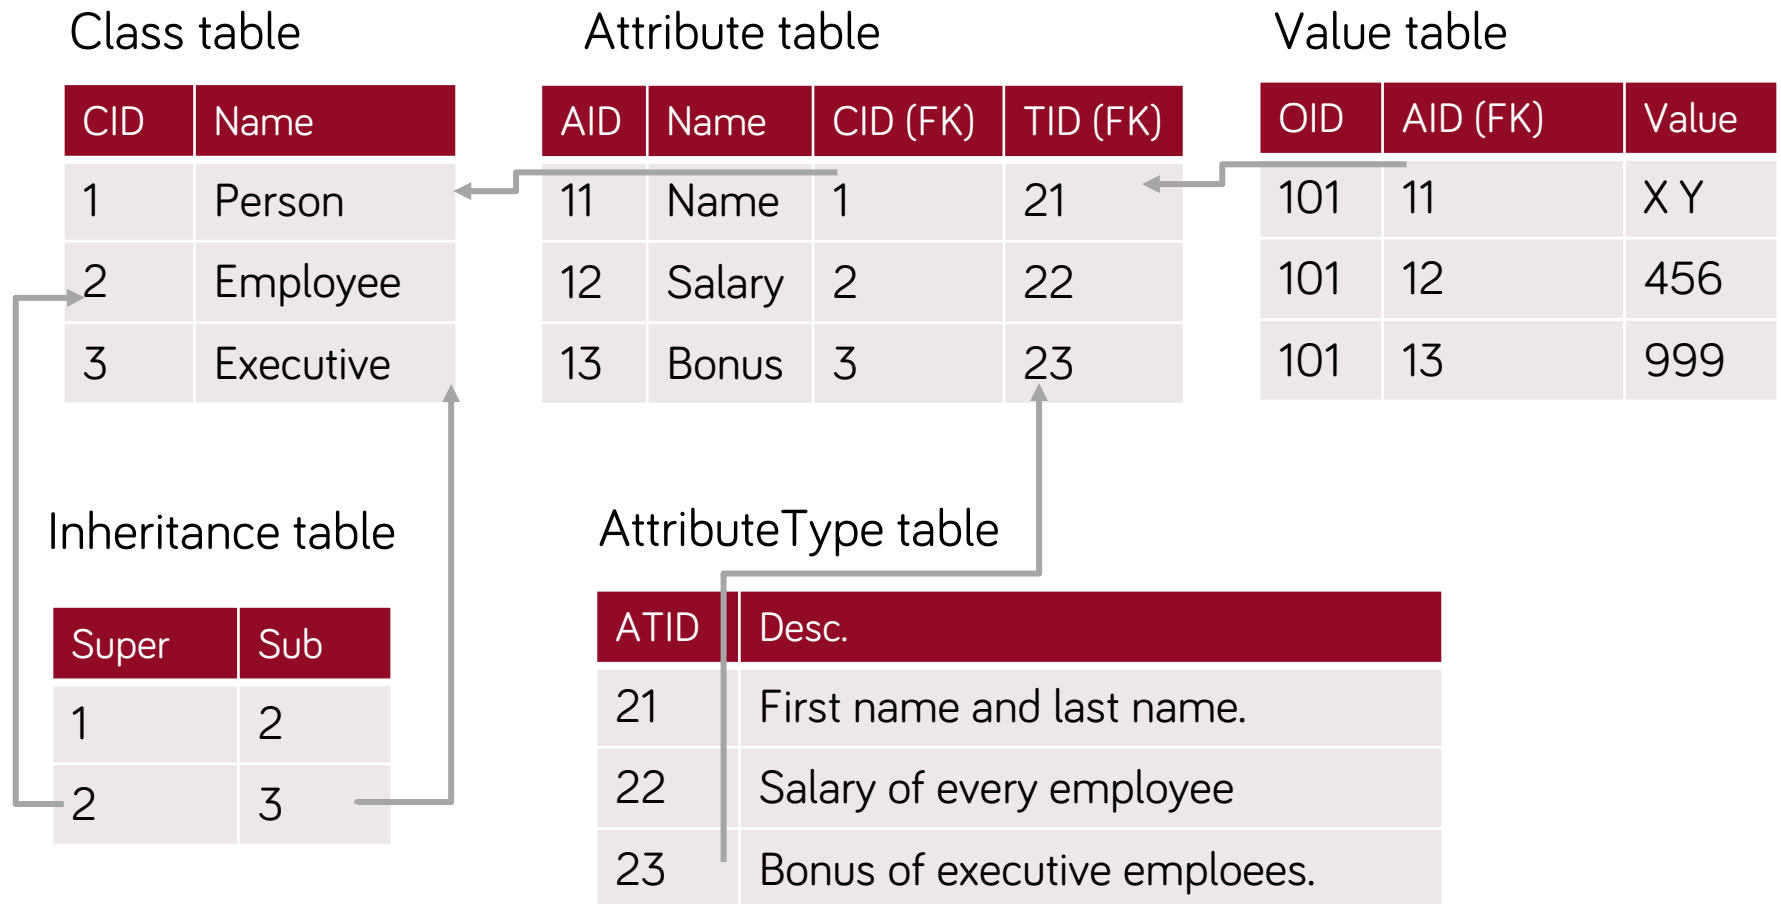
- Meta data driven solution
- General scheme
 - > Can represent any hierarchy
 - > Independent of the specific classes
 - Class hierarchy → Metadata
 - Class instances → Attributes

Mapping into a general scheme - 2



Copyright 2002-2006 Scott W. Ambler

Mapping into a general scheme - 3



Mapping into a general scheme - 4

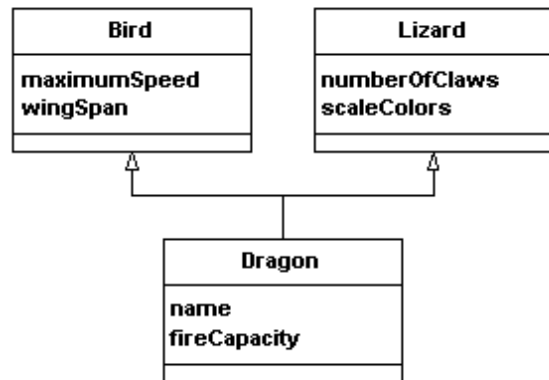
- Advantages
 - > Flexible
 - > Can describe anything
- Drawbacks
 - > Hard to understand at first
 - > Hard to gather all data of an instance
 - > Not efficient for large data sets
- Use when ... - *design time* dimension
 - > Complex applications
 - > Small amount of data
 - > Can change in runtime

Multiple inheritance

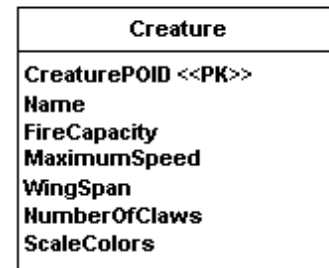
- Modern languages do not support
 - > But C++ → Still supported
- Same solutions are applicable
 - > The general solution includes this case

Multiple inheritance

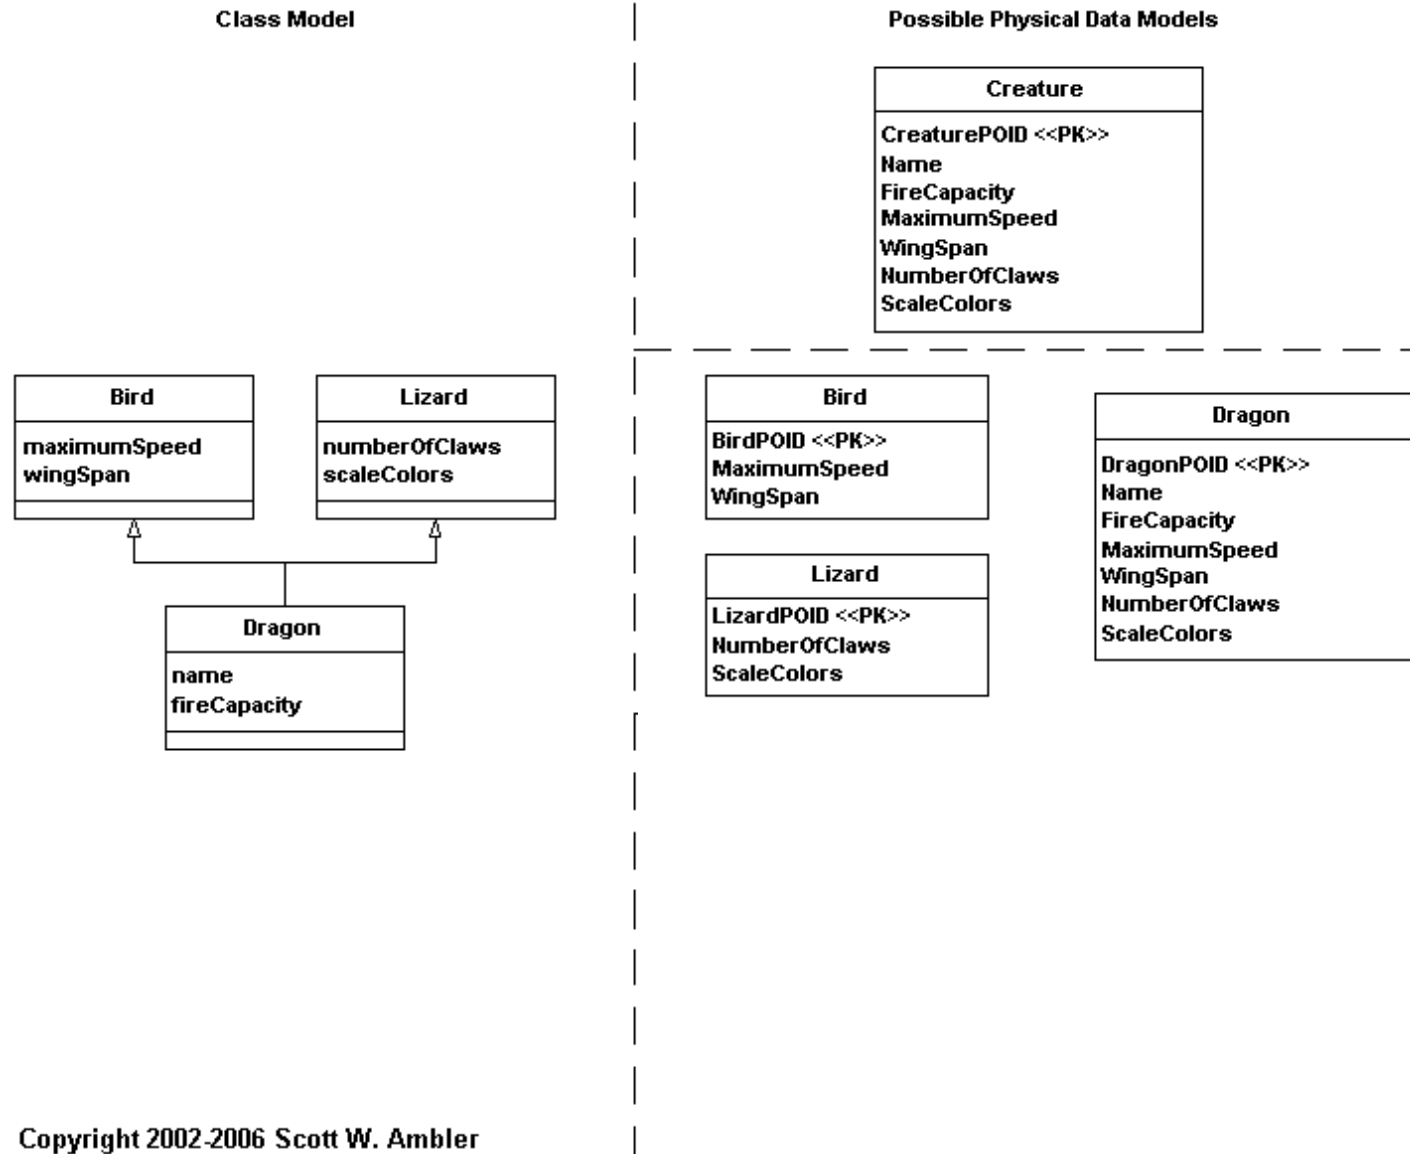
Class Model



Possible Physical Data Models



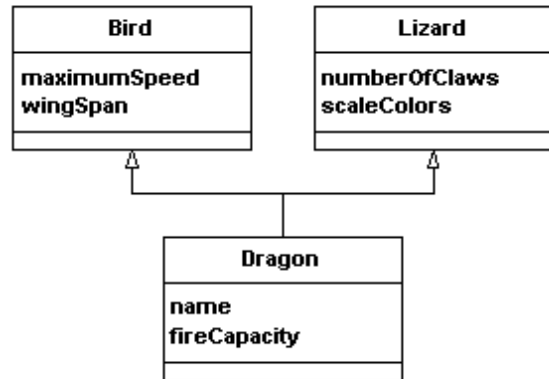
Multiple inheritance



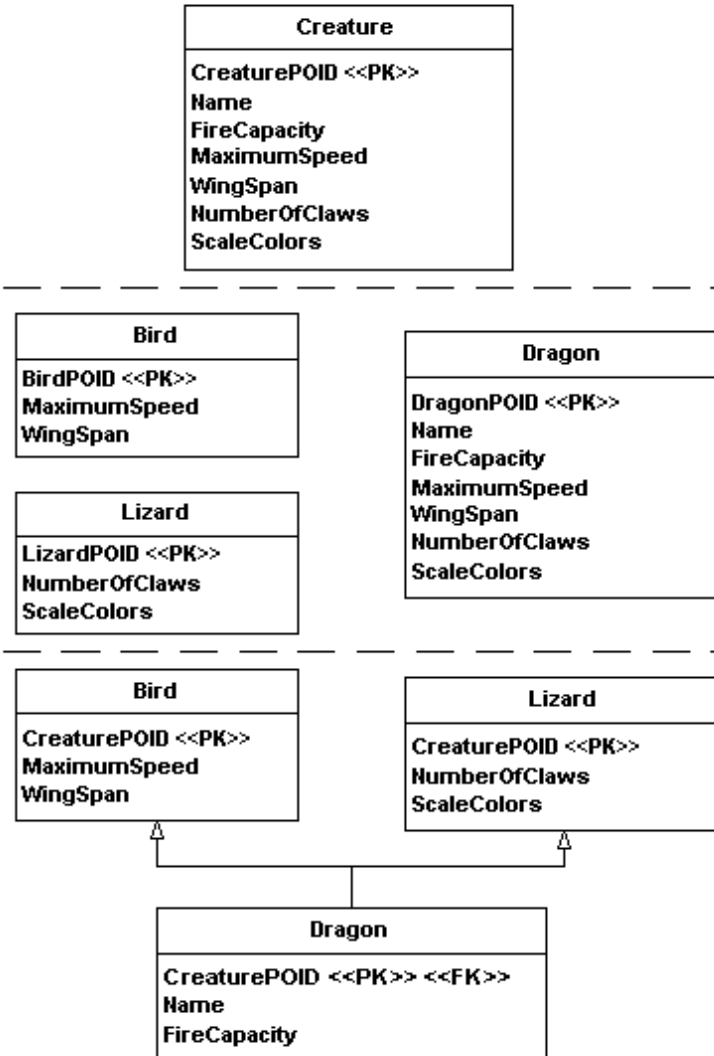
Copyright 2002-2006 Scott W. Ambler

Multiple inheritance

Class Model



Possible Physical Data Models



Copyright 2002-2006 Scott W. Ambler

Which one to use?

- The above mentions design-time aspects are often secondary!
- What queries and actions do we perform and how often?
 - > Filters, orders etc.
- What amount of data do we have?
- What is the distribution of the data?
- You even have to test and measure the alternatives!

Object-relational mapping

Types of relationships

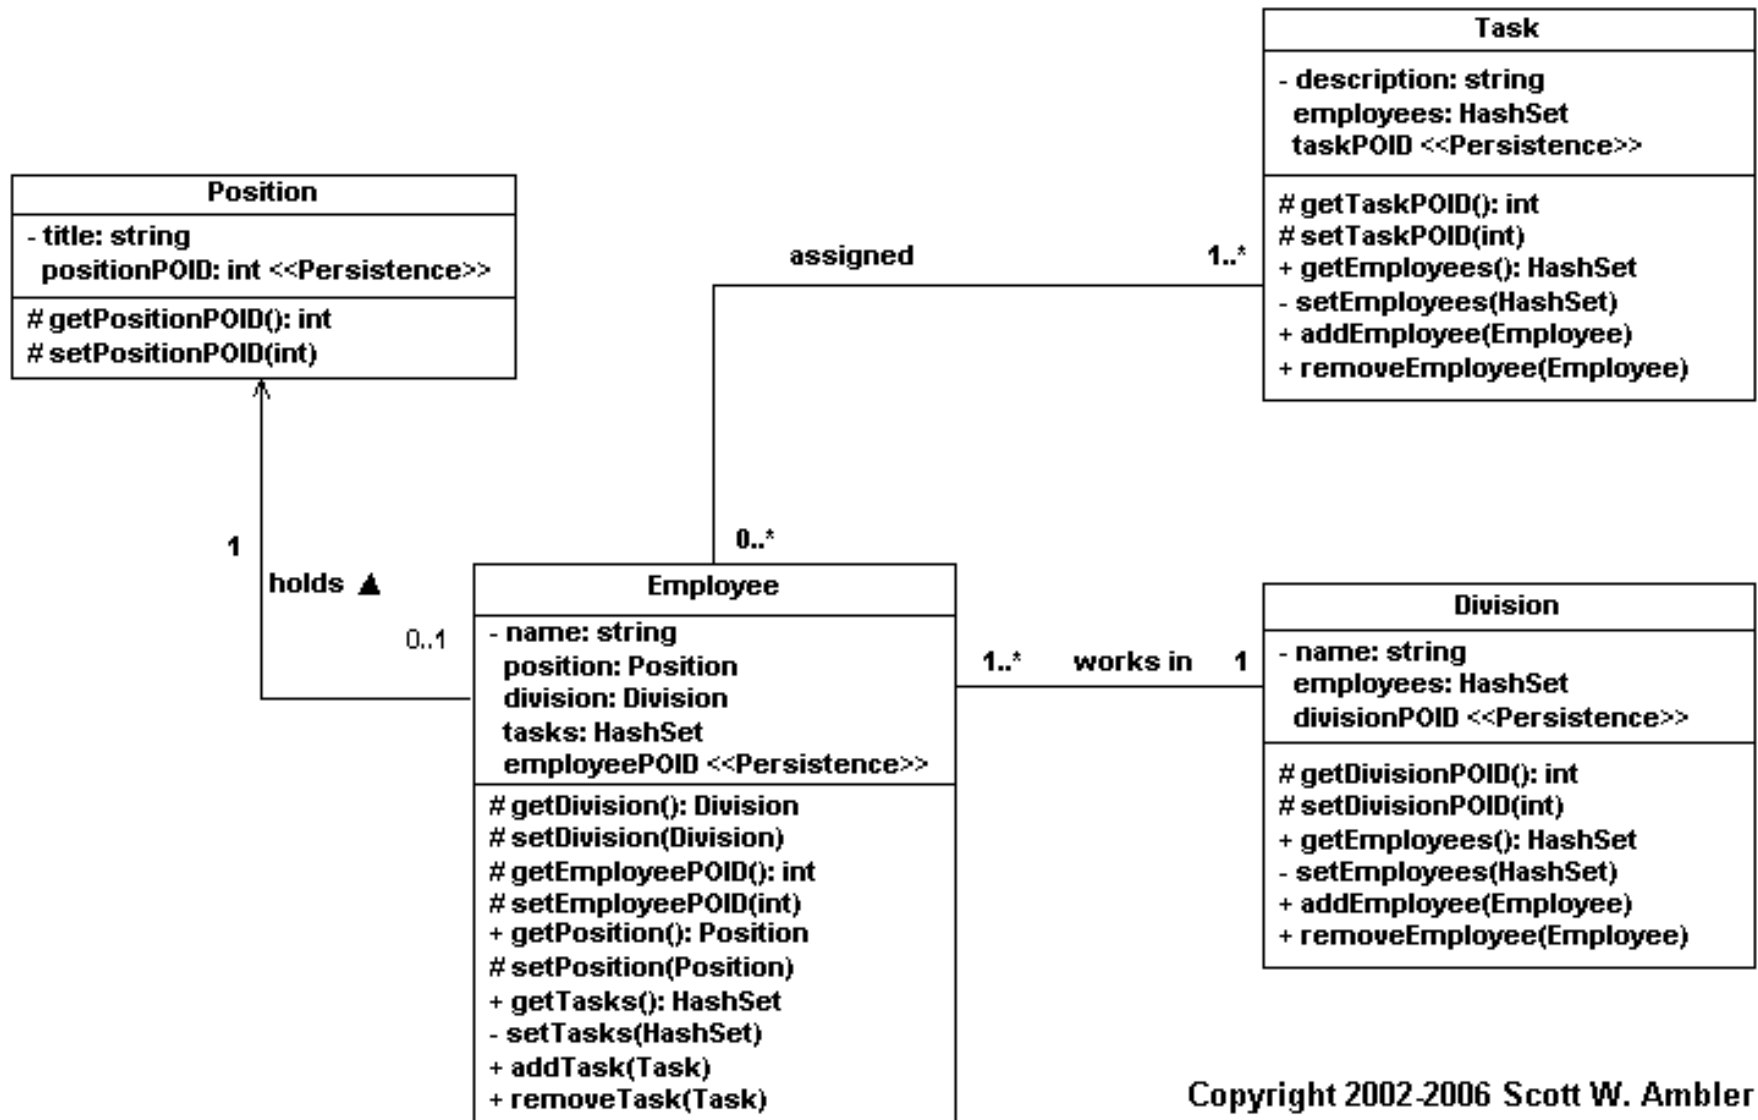
Object Relationships – 1

- Types of relationships
 - > Association
 - > Aggregation
 - > Composition
 - Types
 - > One to one
 - > One to many
 - > Many to many
 - Directed
 - > Oneway
 - > Navigable both ways
- Referential integrity
- Cannot be mapped

Object Relationships – 2

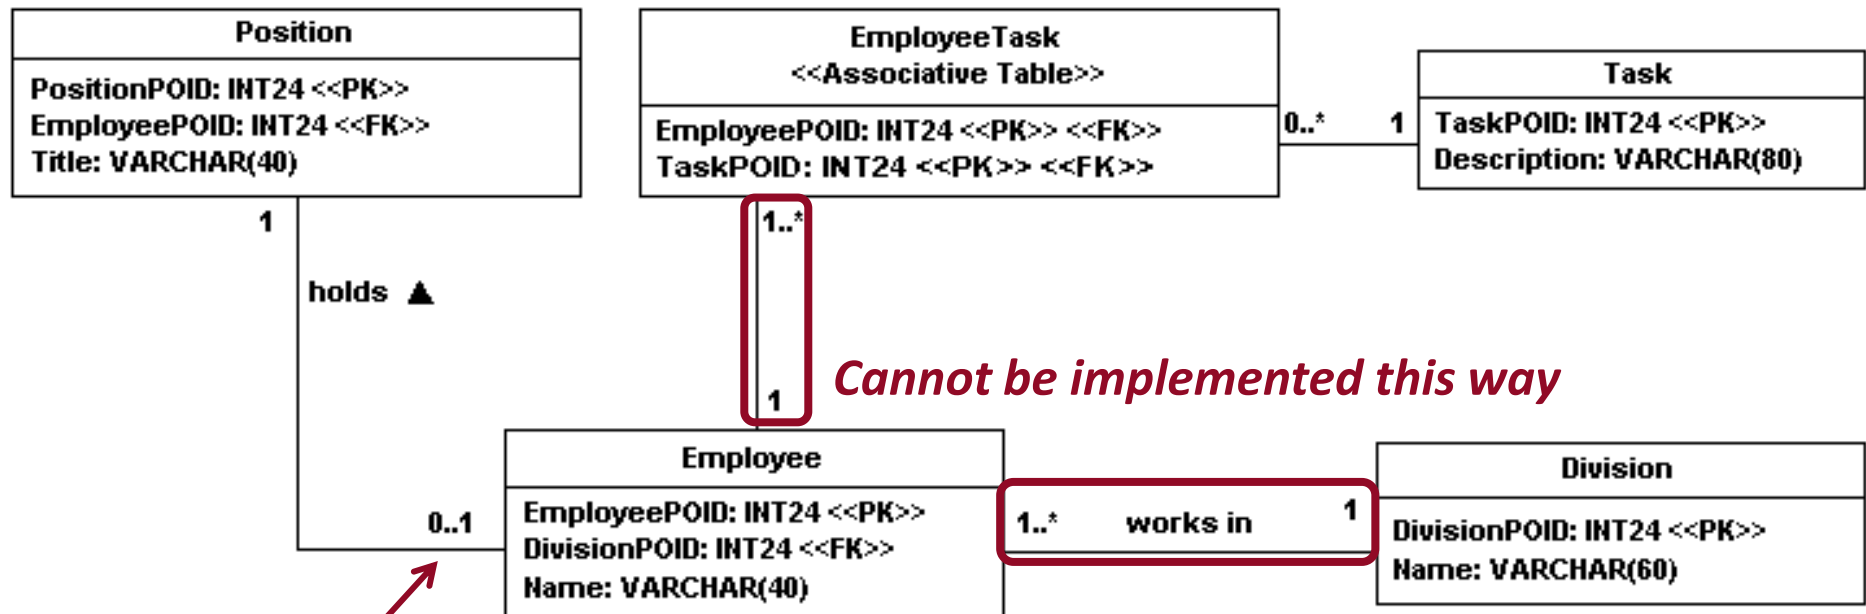
- One to one
 - > Foreign key in one of the tables
 - (May imply one to many connection)
- One to many
 - > Foreign key to the “one”
- Many to many
 - > Cannot be mapped directly
 - > Requires a junction table
- Cardinality
 - > Both ends required - do not map
 - > Hard to start populating the dataset
 - > Cardinalities 0, 1, many

Object Relationships – 3



Copyright 2002-2006 Scott W. Ambler

Object Relationships – 4



Cannot be implemented this way

Could be one to many based on the structure


Copyright 2002-2006 Scott W. Ambler

Recursion – 1

- Also known as reflection
- Both ends of the relationship terminate at the same class
- Like all the other relationships
 - > Little difference in many to many

Recursion – 2


<<Class Model>>



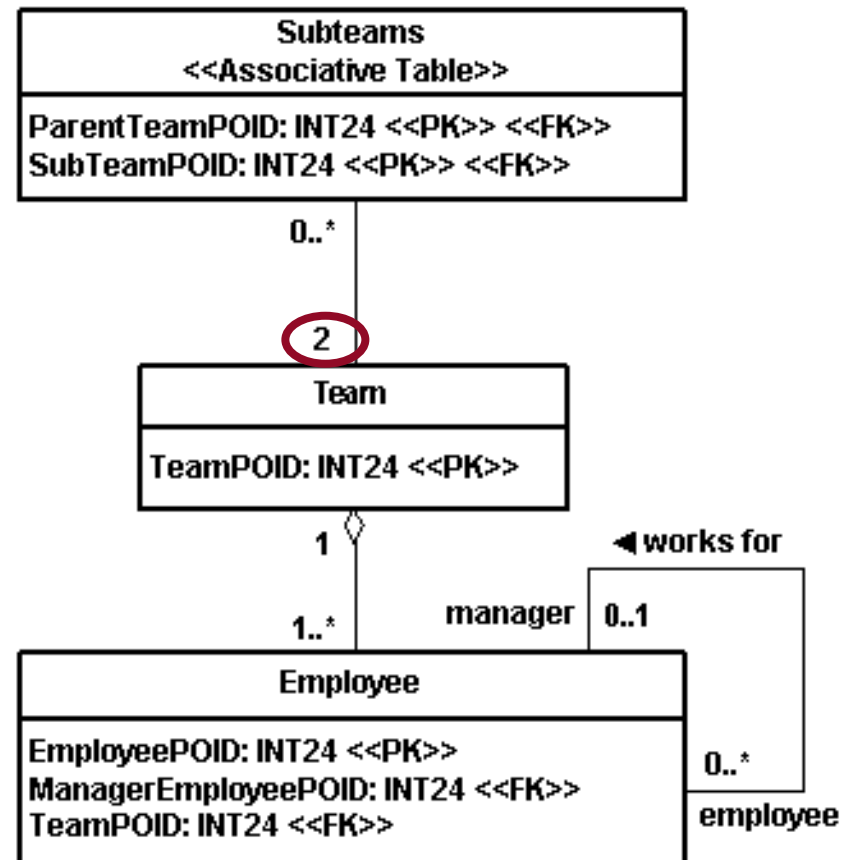
Copyright 2002-2006 Scott W. Ambler

Recursion – 2

<<Class Model>>



<<Physical Data Model>>







Copyright 2002-2006 Scott W. Ambler



Object-relational mapping

Further problems

Ordered collection

Lecture notes and materials

-  Lecture notes, seminar materials and homeworks
-  Introduction and three-tier architecture Uploaded 7/09/20, 13:17
-  Transactions and the MSSQL platform Uploaded 2/09/20, 10:52
-  Microsoft SQL Server programming Uploaded 2/09/20, 10:52



- Attribute the order is based on + read from the DB according to this
- Change in the order may effect
 - > Change in multiple records
 - > Alternative: empty intervals in-between
- Deletion
 - > Not always required to re-index
 - Meaning is sort ordering, not the actual position

Class-level properties - 1

- Specific to the class
- Does not relate to any instance of the class
- Example
 - > Identifier of the next invoice
 - > Fixed percentage of sale today
- Class-level constants

Class-level properties - 2

- Each property into a new table
 - > Fast
 - > Lot of small tables → chaotic data model

InvoiceGlobals table

NextId
123

Sale table

Sale
2%

Class-level properties - 3

- Single table, all properties into specific columns
 - > Fast
 - > Simple (only a single table)
 - > Problematic concurrency (since this is column and not row-based)

Globals table

NextInvoiceId	Sale
123	2%

Class-level properties - 4

- Single table per class
 - > Fast
 - > Lot of small tables → chaotic data model

InvoiceGlobals table

NextId	PaymentDue
123	30

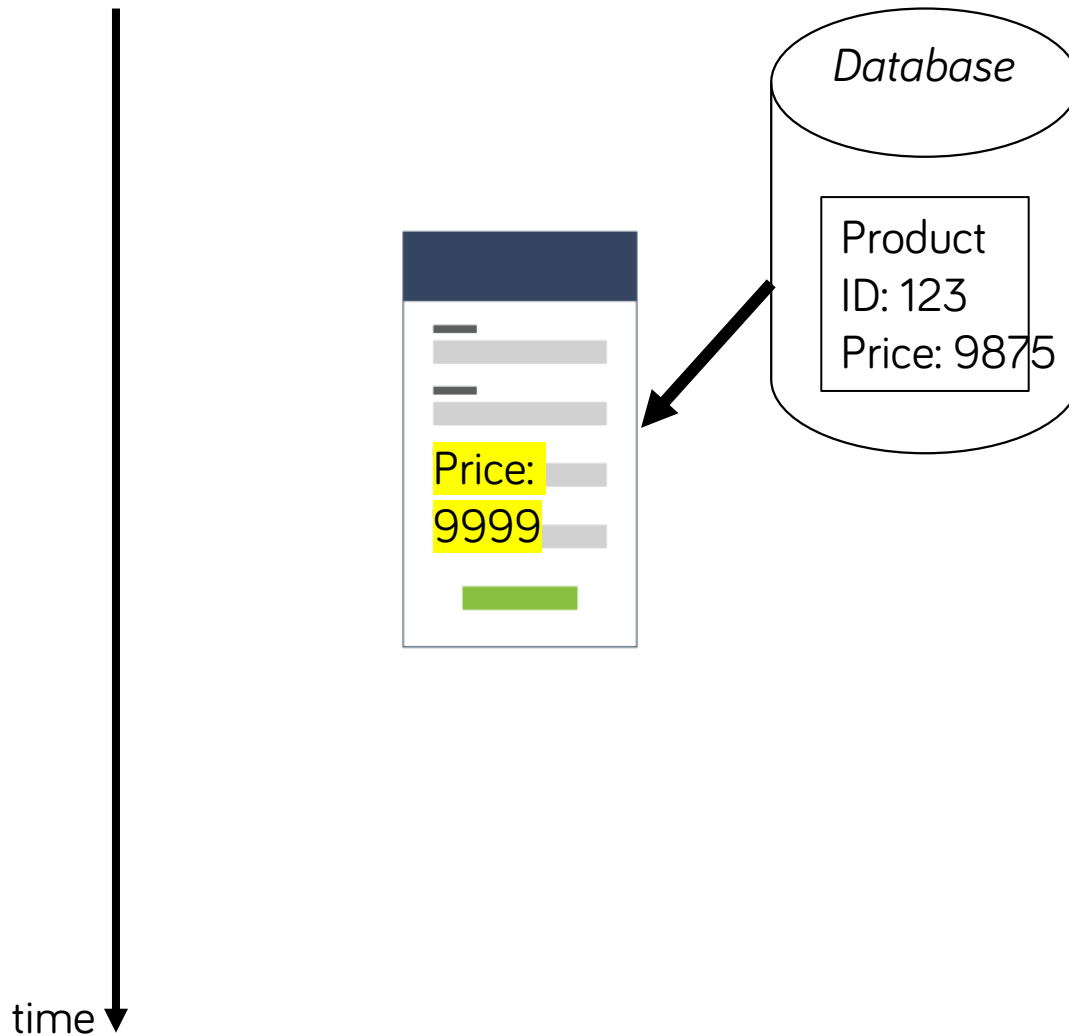
Class-level properties - 5

- General solution
 - > Each property a new row
 - Class
 - Name of the property
 - Value
 - > Have to convert the data
 - > Easy to extend
 - New property → New record

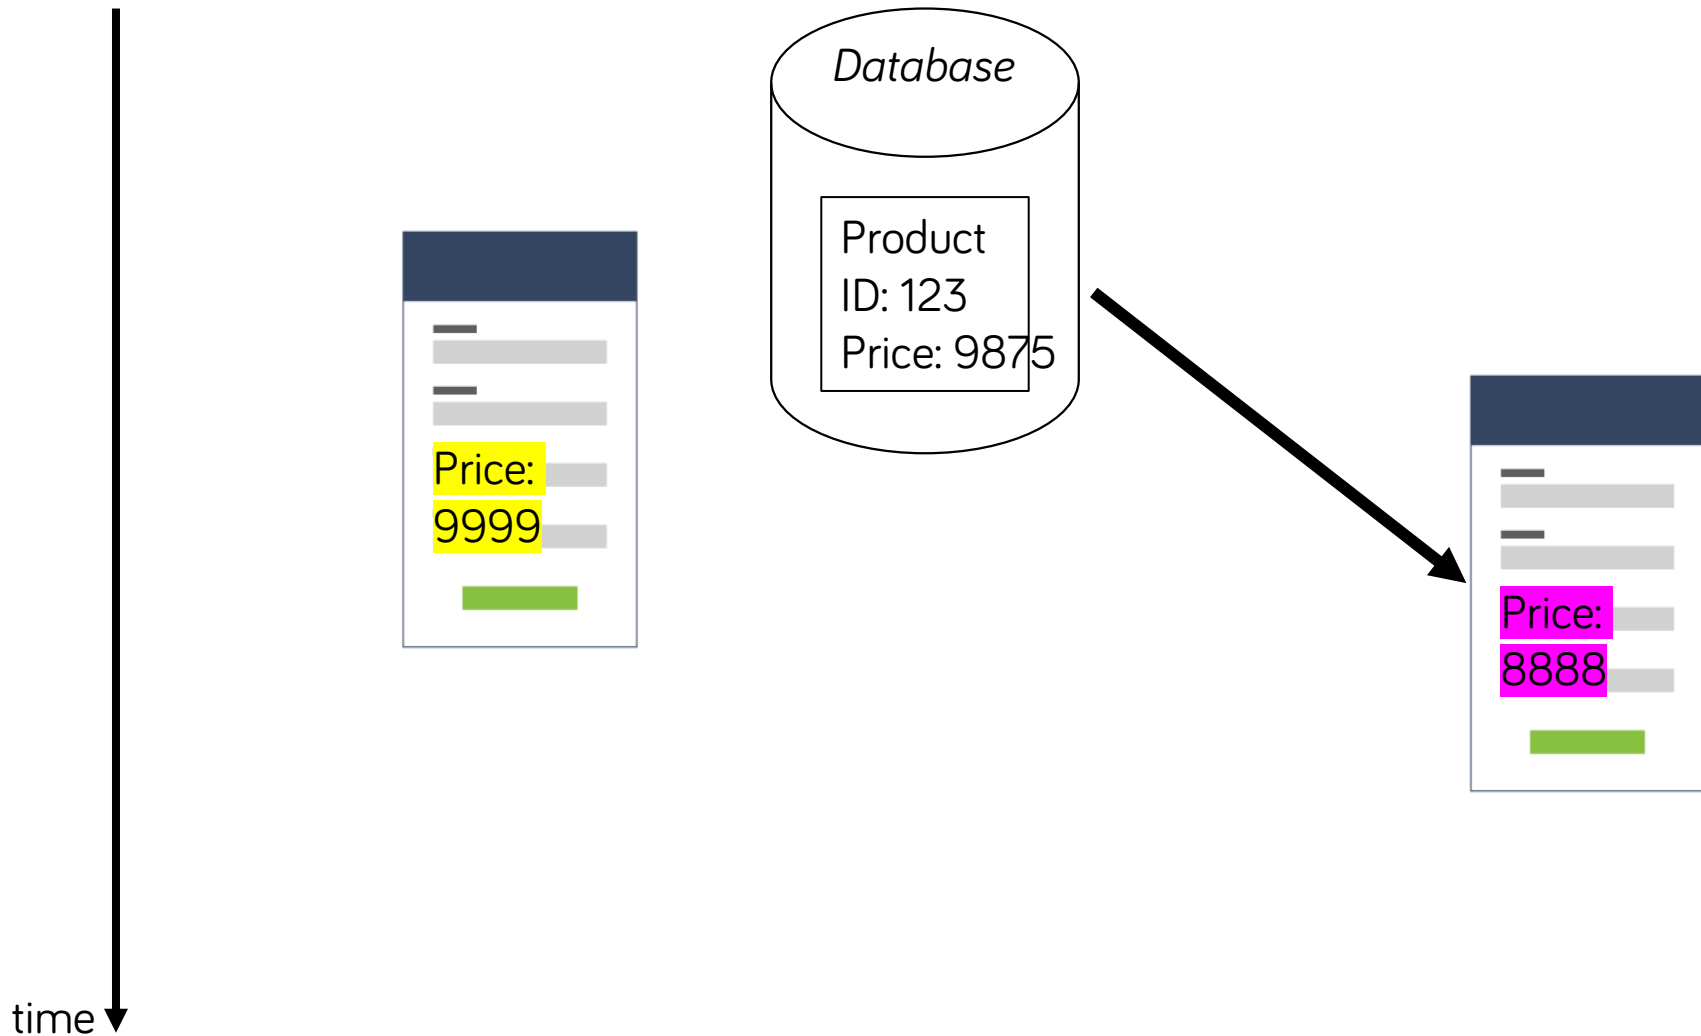
Class	Key	Value
Invoice	NextId	123
Invoice	Sale	2

Concurrency handling

Concurrency handling: problem



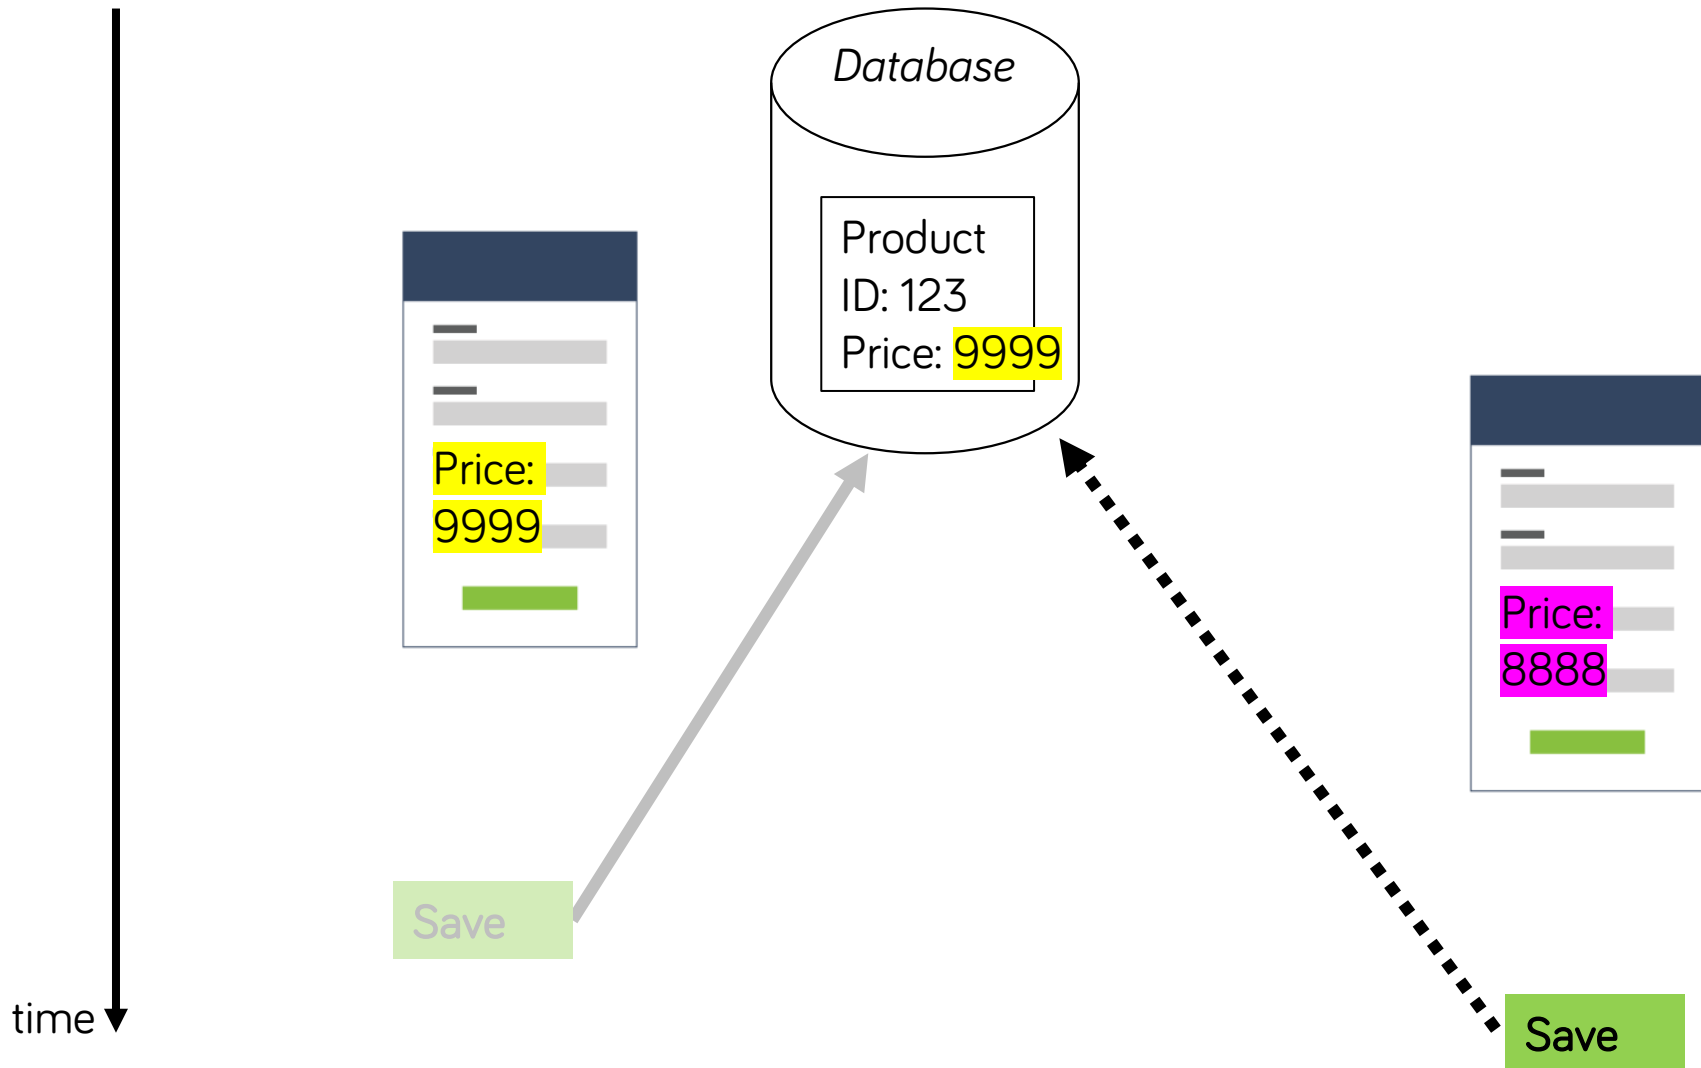
Concurrency handling: problem



Concurrency handling: problem



Concurrency handling: problem



Concurrency handling

- None – seldom a good solution 😞
- Pessimistic concurrency handling
 - > Someone else is going to modify the same data
 - > Get exclusive access to the data
 - Locking + Transactions
- Optimistic
 - > No one will modify
 - > But we must detect clash

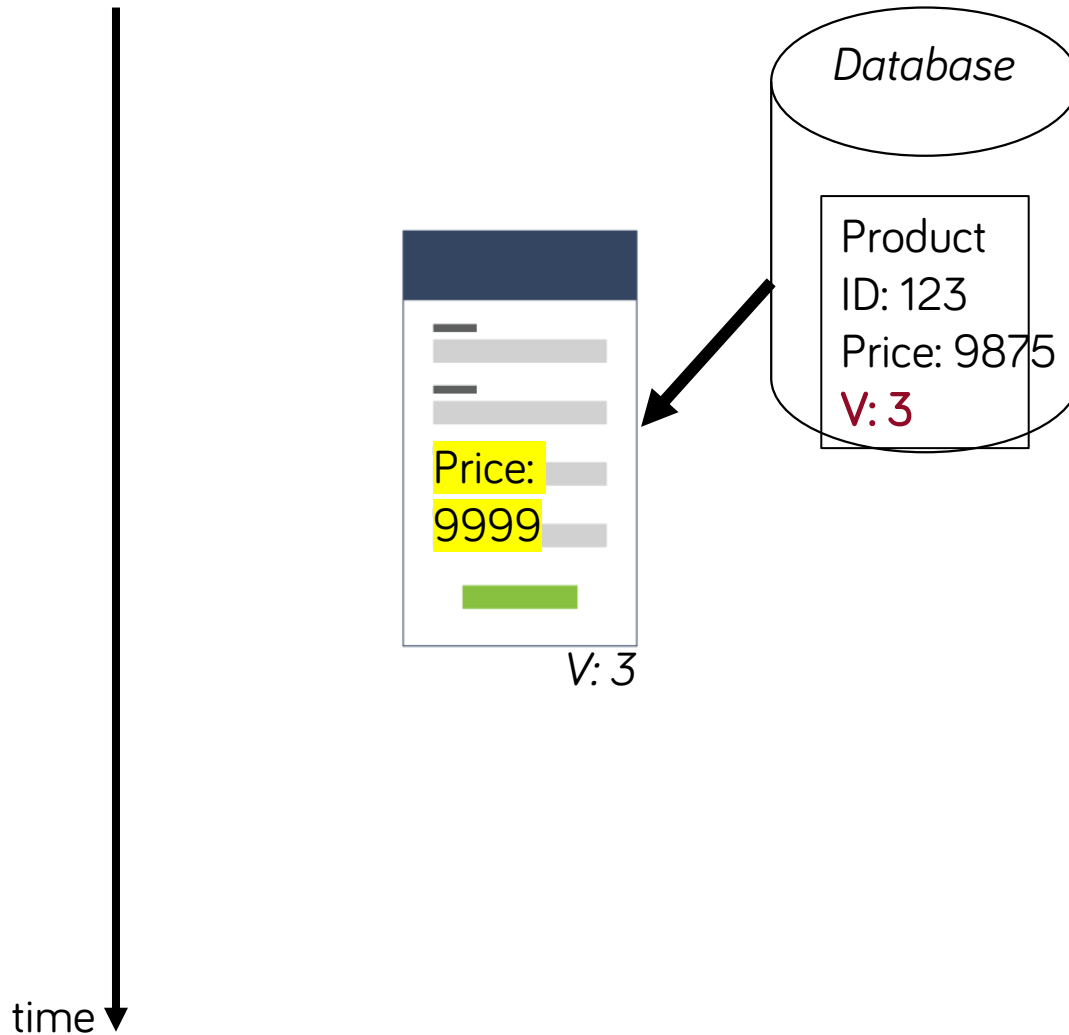
Pessimistic concurrency handling

- When a connection with the database is maintained
 - > Locking individual records
 - > Database transactions do this
- If the connection to the database has been closed already (UI)
 - > Typically happens in the Business Logic Layer
 - > Exclusion handled manually in the Business Logic Layer
 - > Scalability + SPOF
 - > Handling closed session

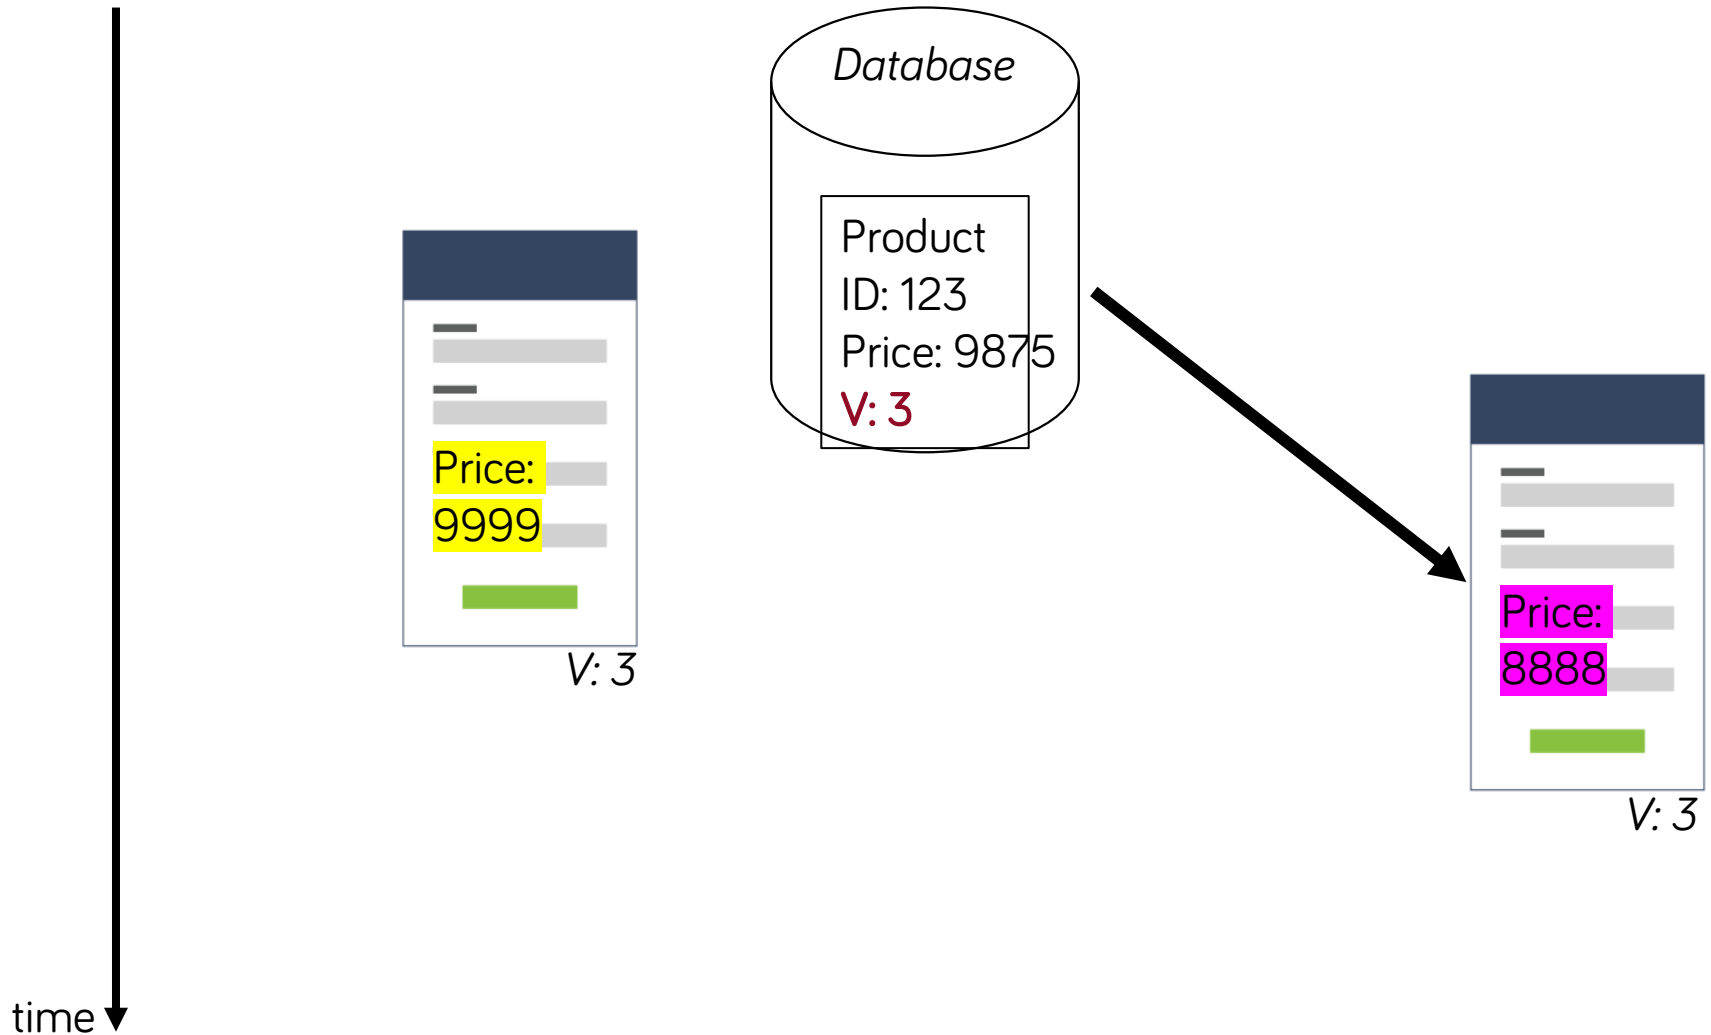
Optimistic concurrency handling

- Must check changes before committing them
- Compare record contents
 - > Record versioning
 - Logical counter
 - Requires database scheme change
 - > Compare the contents
 - Preserve the state before the changes
 - Additional information within the business entities
- The updating SQL command must be written in a specific way
 - > Data access layer provides this

Optimistic concurrency handling



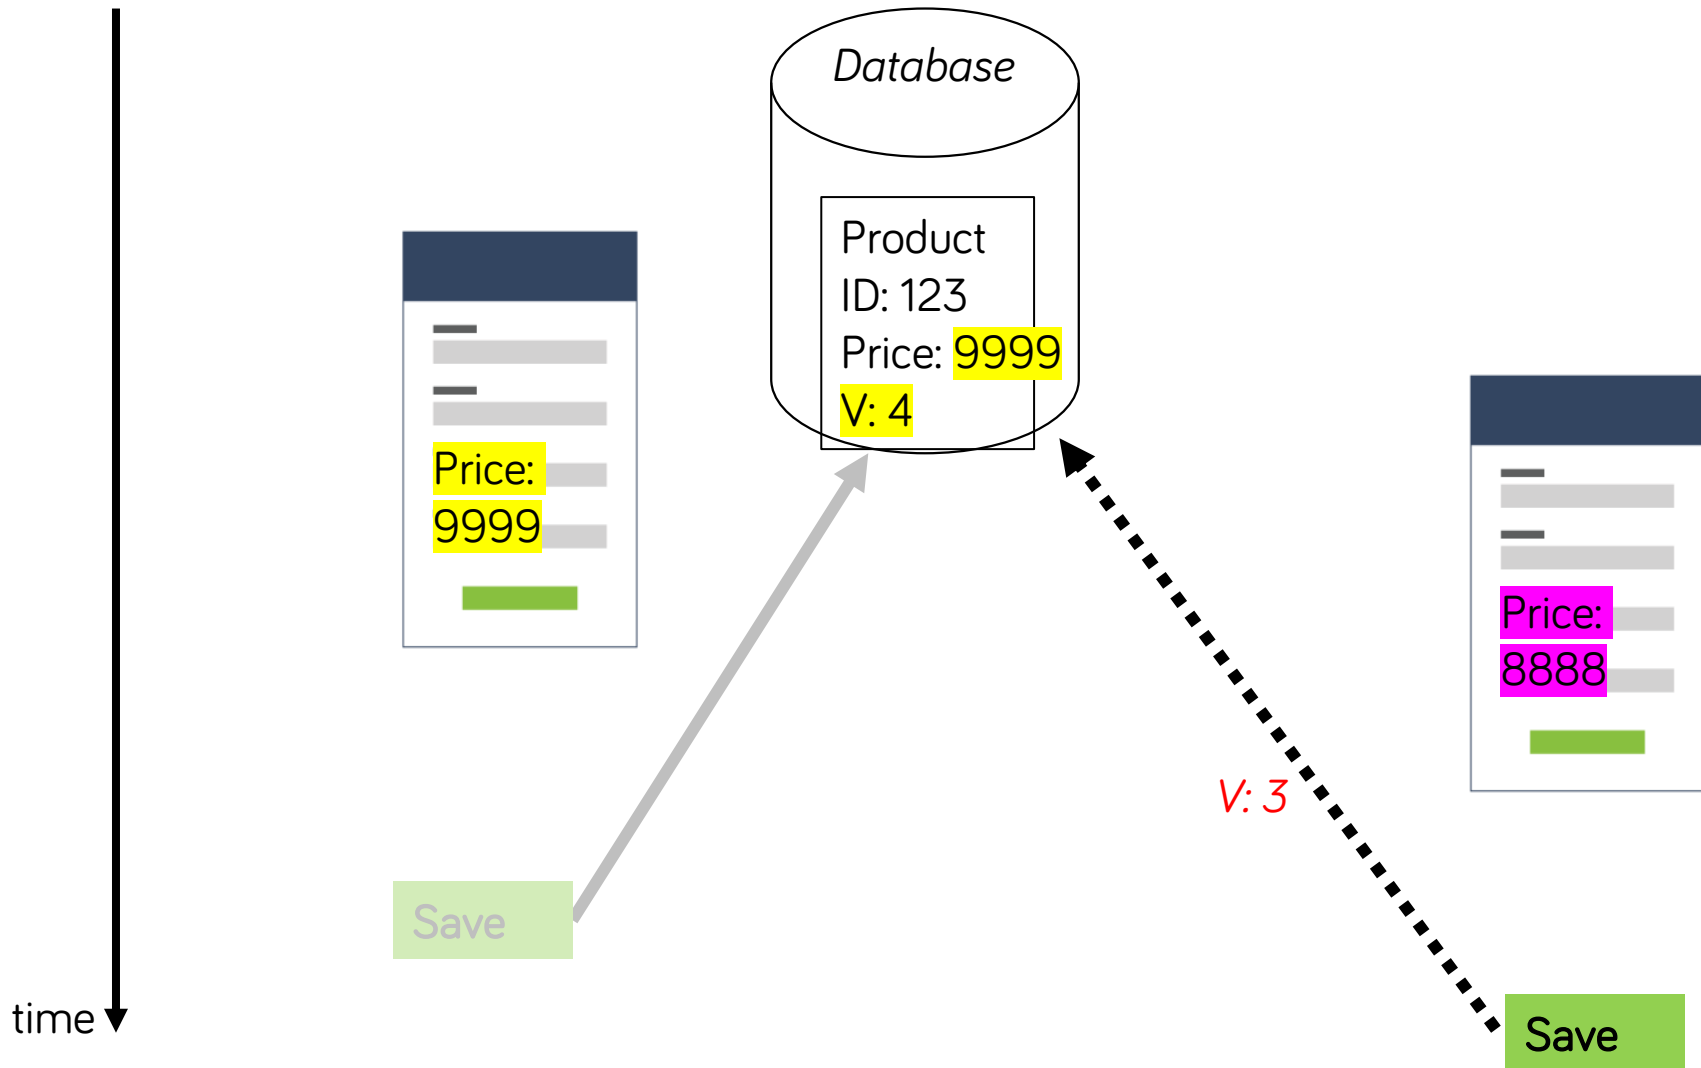
Optimistic concurrency handling



Optimistic concurrency handling



Optimistic concurrency handling



Optimistic concurrency handling: resolution

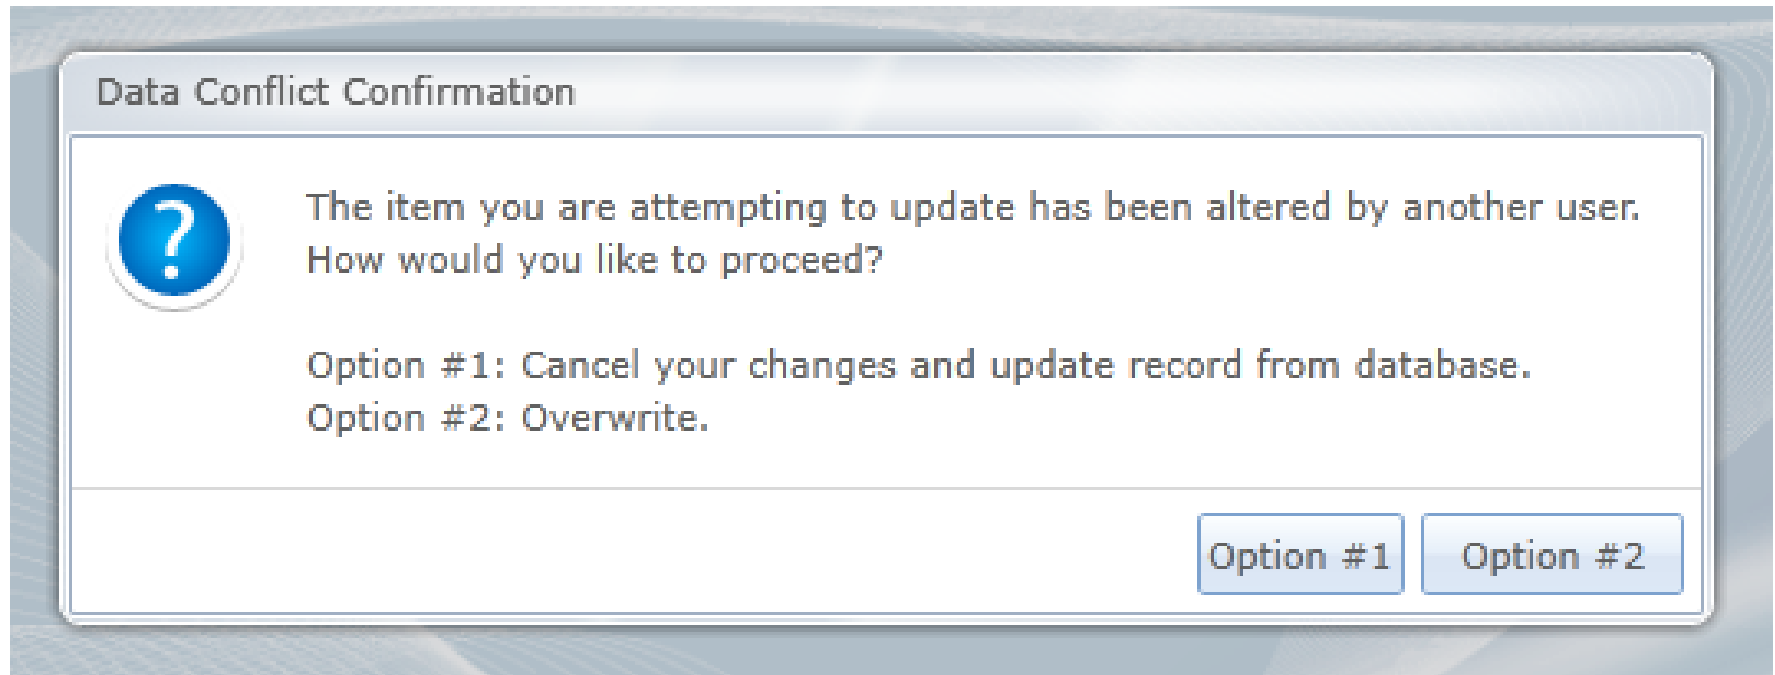


Image source: <http://reptonprojects.com/index.php/application/concurrency-control>

Detection strategies

1. The first writer wins 🤪
2. Last writer wins 😐
3. We leave it to the user 🤖
4. We compare the set of modified fields 🤔
 - > Consistency must be maintained!
5. In the vast majority of cases, the changes do not only affect a single table, but several records! 🤪 🤪


Detection -> API

- In the case of a multi-layered application, the data is transferred to the UI layer
- The database connection is closed
- In case of a modification request, it is necessary to know **from somewhere** what the original status / version of the records were
- The server is stateless -> the UI must store information about the original record
 - > -> the API changes

Design patterns in the data-access layer

Repository


Multiple data storage solutions



Solution

- Place the parts to be replaced behind an interface!
- Loose coupling
- It is clearly visible what needs to be rewritten
- It can be tested
- Tasks can be better divided between teams/devs

Multiple data storage solutions



Advantages of data access layer separation

- Data access layer replacement
 - > Linq to EF, nHibernate, ...
- Because of unit test, for mocking
- Database refactoring
- You need to switch to another technology because it generates better SQL for the database engine
- It is necessary to switch to another database that is not or poorly supported by the current provider
- ...

Is it necessary to replace it?

- YAGNI: You aren't gonna need it
 - > "do the simplest thing that could possibly work"
- KISS: Keep it simple and stupid
- **Agile:** don't plan and implement the system based on *possible future* requirements!
 - > Customer is paying when he wants it...
- **Waterfall:** think hard about whether this kind of flexibility will be needed during the application's lifetime
 - > If so, can the costs be passed on?

Alternative solutions

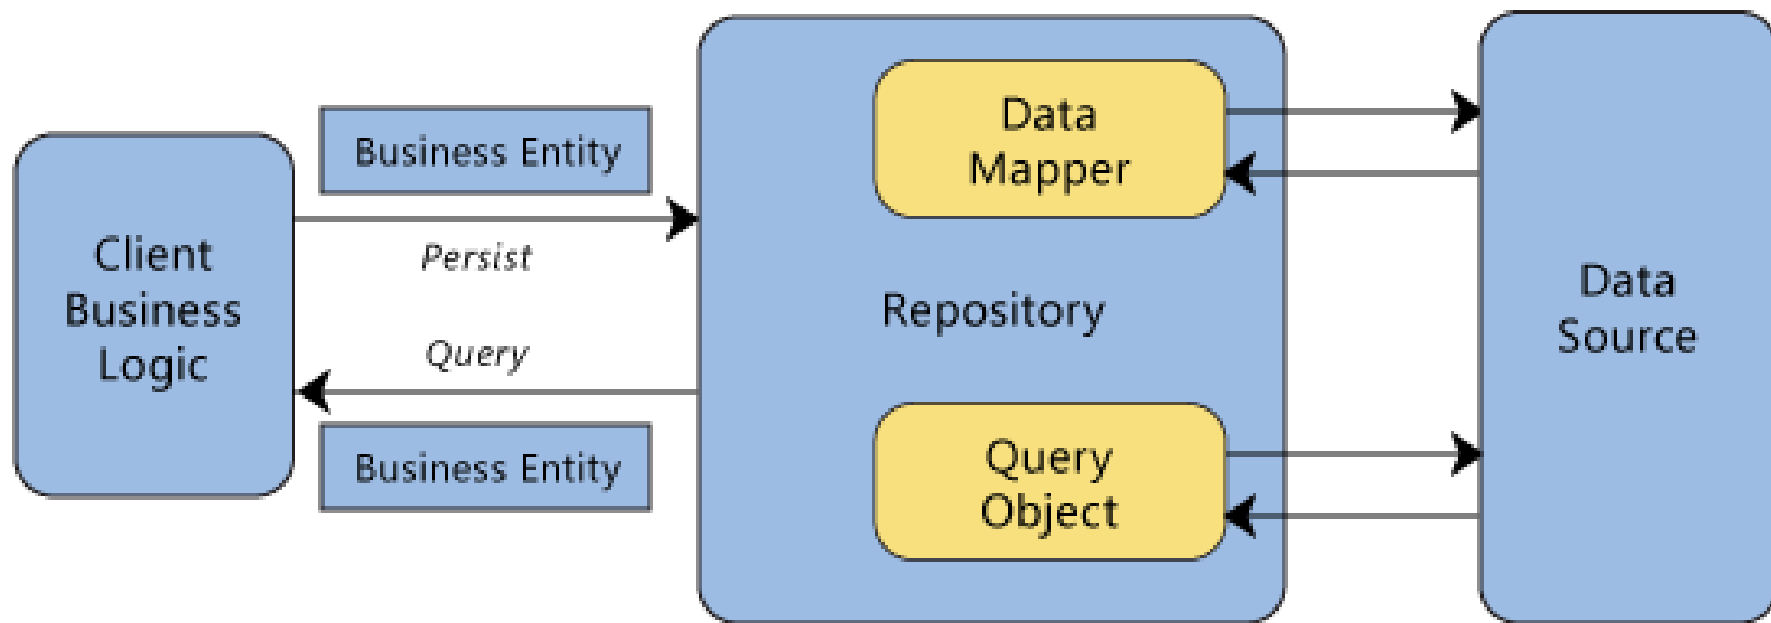
- What is the exact reason why this might be necessary?
- Unit testing?
 - Search in-memory database, SQLite, ...
- Replacing the database?
 - Look for paid components, see if there is a provider
- To what extent should the data layer be hidden from the business logic?

Separation of data (access) layer

- Separation: typical interface design
- The services and functions of the lower layer are explicitly exposed on the interface
- A widely used solution:

Repository pattern

Data access abstraction: repository pattern



Kép forrása: [https://docs.microsoft.com/en-us/previous-versions/msp-n-p/ff649690\(v=pandp.10\)](https://docs.microsoft.com/en-us/previous-versions/msp-n-p/ff649690(v=pandp.10))

Repository sample

API



Domain



Repository



Data access

```
public User GetById(int userId)
```

```
public class User
{
    public int UserId;
    public string FirstName;
    public string LastName;
}
```

```
public class UserRepository : IUserRepository
{
    public User GetById(int userId)
    {
        //DB Call here...
    }
}
```

Abstracting data access



Repository: class

- Multiple repositories
 - > Per entity or entity group
 - > Product and category; Orders; ...
- One repository: one class
 - > It is typically used with an interface
 - Testing, mocking (see later)
 - > CRUD operations
 - > Works with business entities
- All database-specific details are enclosed here
 - > Technology (EF Core) and platform dependent (MS SQL)

Repository sample: CRUD

```
class ProductRepository : IProductRepo
{
    List<Product> List() { ... }
    Product FindById(int Id) { ... }
    void Add(Product entity) { ... }
    void Delete(Product entity) { ... }
    void Update(Product entity) { ... }
}
```

Repository sample: operations

```
class ProductRepository : IProductRepo
{
    void AddProductToCategory(
        Product p, Category c) {...}

    void StopSellingProduct(
        Product p) {...}
}
```

Repository Pattern properties

- Abstraction above the data layer
- A specific location where the way data is accessed can be changed
- Location of datatables
- It can easily be changed to another implementation
- It hides the details
- Several alternative implementations

Repository Pattern

- Business logic independent of data layer
- Separation of domain logic and data access
- Data access/database replacement
- Unified management of multiple data sources
- Change of data storage paradigm (e.g. nosql, lucene, etc.)
- Support for unit tests
- Better parallelization of development
- Reduction of duplicate queries
- Better control over queries
- Object oriented API
- Uniform rules
- Caching
- Another layer - more work
- It reduces the power of ORM technologies
- Risk: Data storage API still leaks into business logic

ORM and the repository

- Different goals
- Repository
 - > Abstraction and consolidation of all functions related to storage
 - > Architectural pattern
- ORM
 - > Abstraction of supported relational database access
 - > Used behind Repository as an "implementation detail"

Repository functions in practice

- Add and remove entities
- The interface uses tables and collections
- Transaction management/closing is not here
- Queries according to the desired criteria

Query functions

```
public Customer[] WithSurname(string surname)
{
    :
    :
}
```

Generalized queries

```
public Customer[] Find(ICustomerSpecification spec)
{
    ...
}
```

```
public interface ICustomerSpecification
{
    bool IsSatisfiedBy(Customer c);
}
```

- Part of the domain model
- Composable (complex expressions)
- Must work in-memory or generating SQL

Generic repository pattern

```
interface IRepository<T>
{
    T[] Find(ISpecification<T> specification);
}
```

Rhino Commons

{ } Rhino.Commons

IRepository

Get(object id):T

FutureGet(object id):FutureValue<T>

FutureLoad(object id):FutureValue<T>

Load(object id):T

Delete(T entity):void

DeleteAll():void

DeleteAll(DetachedCriteria where):void

Save(T entity):T

SaveOrUpdate(T entity):T

SaveOrUpdateCopy(T entity):T

Update(T entity):void

FindAll(Order order, params ICriterion[] criteria):ICollection<T>

FindAll(DetachedCriteria criteria, params Order[] orders):ICollection<T>

FindAll(DetachedCriteria criteria, int firstResult, int maxResults, params Order

FindAll(Order[] orders, params ICriterion[] criteria):ICollection<T>

FindAll(params ICriterion[] criteria):ICollection<T>

Sharp Architecture

```
IRepository<T> (in SharpArch.Core.PersistenceSupport)
IRepositoryWithTypedId<T,IdT> (in SharpArch.Core.PersistenceSupport)
    Get(IdT id):T
    GetAll():List<T>
    FindAll(IDictionary<string,object> propertyValuePairs):List<T>
    FindOne(IDictionary<string,object> propertyValuePairs):T
    SaveOrUpdate(T entity):T
    Delete(T entity):void
    DbContext:IDbContext
```

Fluent NHibernate

```
IRepository (in FluentNHibernate.Framework)  
... Find<T>(long id):T  
... Delete<T>(T target):void  
... Query<T>(Expression<Func<T,bool>> where):T[]  
... FindBy<T,U>(Expression<Func<T,U>> expression, U search):T  
... FindBy<T>(Expression<Func<T,bool>> where):T  
... Save<T>(T target):void
```

Is IQueryable usable here?

Using IQueryable<T> GetAll<T>() method in BLL:

```
public Customer[] GetCustomersPaged(int pageNumber, int pageSize)
{
    var customers = _customerRepository.GetAll()
        .Where(c => c.IsDeleted == false)
        .Skip(pageNumber * pageSize)
        .Take(pageSize);
}
```

```
interface IRepository<T>
{
    IQueryable<T> GetAll();
}
```

- The SQL executed outside the repository!
- Query technology can also infiltrate domain areas of the application ☹️
- Called „Leaky abstraction”

About transactions

- It is not the responsibility of the repository
- Unit of Work pattern
 - > NHibernate session
 - > EntityFramework DataContext
- Establishing transaction boundaries is the responsibility of business logic

Criticism of the repository pattern

- **"Redundant"**: no additional abstraction layer is needed
- It **complicates** the code, since "every" query must be written explicitly, LINQ cannot be used
- If we use IQueryable, we lose the benefits of the pattern because the query is ultimately generated by the BLL and we are tied to IQueryable
- "Return" to the world of stored procedures
- **Testing**: generally not easier to mock

Advantages of the repository pattern

- **Domain-driven:** provides a domain-specific view of the database - as opposed to queries formulated in LINQ, which are general
- **Hide/Aggregation:** hides the complexity of related entities (/tables) and gives a clearer view
- **Testing:** it's easier to mock for a small project

Guideline

- What is the goal?
 - > Testing, maintenance, responsibilities, ...
 - > Don't get caught up in the names (BLL, DAL, repo, ...)!
- What is business logic? Where is?
 1. In queries
 - > With the repository pattern, the BLL can be almost "empty", the repository contains a lot of logic
 - > Thus, the repo must be tested when we are testing!
 2. Other logic, on data in-memory?
 - > These codes should be separated and tested there!
 - > They can get the data from the repo or EF...
 3. Other validation logic?
 - > You can relocate to separate classes, etc.

EF and Repository

- Let's refine the repository pattern as we saw earlier
 - > Not just an ID-based query...
 - > Any conditions, filters can be specified
 - > Add group, join, query, ...
 - > Let's enable transactions on the interface
- But this is the EF itself!
 - > A very generic query interface
 - > The implementation is in the database provider
 - > DbContext is just a UoW implementation
- In addition to EF, do not use a repository because it already is a repository implementation
- The repository pattern is good in itself, so otherwise use it (like in Java for example)!