

Data-driven systems

MongoDB



Automatizálási és
Alkalmazott
Informatikai Tanszék

Contents

- „NoSQL”
- Basic concept
 - > System architecture
 - > Collection, document
 - > Data representation
- CRUD operations, query
- Usage in .NET
- “Schema” design

NoSQL: motivation

Small database, simple schema

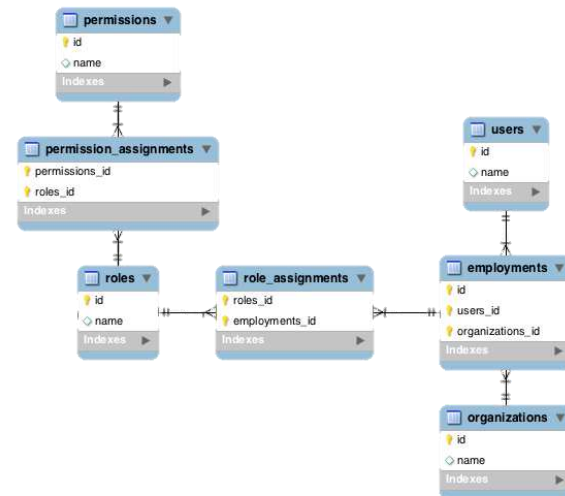



Image source: <https://i.stack.imgur.com/cAwJ2.png>

NoSQL: motivation

Application evolves -> schema gets complicated -> database grows

DB

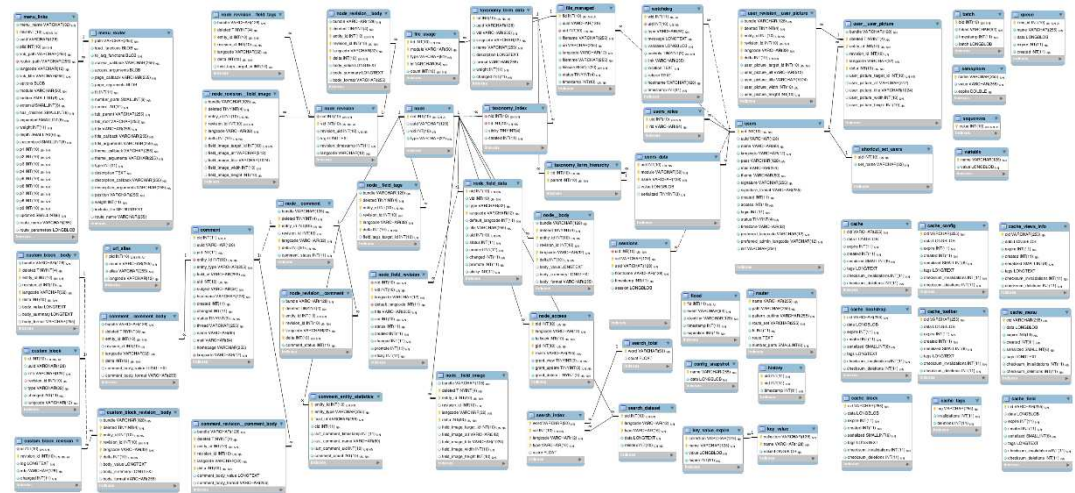


Image source:

<https://www.activequerybuilder.com/blog/understanding-a-complex-database-structure/>

NoSQL: motivation

Real life example

- Backend application developed for 10+ years
- 95+ tables
- Some tables having 50+ columns
- 5500 lines „create database.sql”



NoSQL: motivation

- Drawbacks when using relational databases
 - > Schema is constantly changing
 - > Data migration
 - > Performance issues
 - Scalability is an issue due to strong consistency

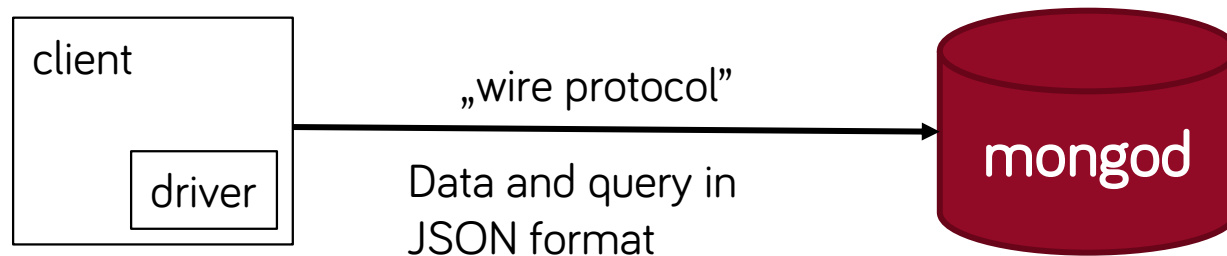
→ Let us get rid of the strict schema: NoSQL

NoSQL: the name is misleading, has nothing to do with SQL

Basic concepts of MongoDB

Architecture

- System architecture



- Logical structure
 - > Cluster
 - Server
 - Database
 - Collection
 - Document

Document

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

Document

- JSON/BSON
- Unit of storage
- Contains key-value pairs
- Value: string, number, null, date, binary, array, nested object, ...
- Key
 - > Free text
 - > Unique name
 - > Cannot begin with character \$
 - > _id implicit field
 - > caseSensitive
- In the object-oriented world: an object
- Size limitation: 16MB

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

Collection

- Analogous to the table of a relational database
- There is no scheme, hence no need to define it
- Place of “similar” documents
- Indices are defined on collections
- Since there is no schema, there are no integrity requirements either

Collection

```
{  
  name: "Sue",  
  age: 26  
}  
{  
  name: "John",  
  age: 32,  
  title: "CEO"  
}  
{  
  name: "Ronald McDonald",  
  age: "born in 1955"  
}
```

Database

- Same purpose as in relational databases
- Umbrella over all data of an application
- Access rights is granted on the database level
- Name is case sensitive, typically used as lowercase

Relational schema / MongoDB

- Table
- Record
- Column
 - > Scalar
- Integrity criteria
- Key
- Foreign keys
- Join
- Transactions
- Collection
- Document
- Field
 - > Scalar & compound (nested)
- ~~Integrity criteria~~
- ObjectId, unique index
- Reference based on _id
- Nesting, arrays
- “*Transactions*”

Key

- In every document: `_id`
- There are no keys otherwise
 - > Cannot define a key
- `ObjectId("507f191e810c19729de860ea")`
 - > Client driver or server generates it (not the application)
- 12 bytes
 - > 4 bytes timestamp (seconds precision)
 - > 5 bytes random (cluster node + mongod process id)
 - > 3 bytes counter
- ~ globally unique

Key

- What the key is used for
 - > Unambiguous identification → `_id`
 - > Guarantee uniqueness → index
- There is no compound key
 - > Compound indices instead

Reference to another document

user document

```
{
  _id: <ObjectId1>,
  username: "123xyz"
}
```

contact document

```
{
  _id: <ObjectId2>,
  user_id: <ObjectId1>,
  phone: "123-456-7890",
  email: "xyz@example.com"
}
```

access document

```
{
  _id: <ObjectId3>,
  user_id: <ObjectId1>,
  level: 5,
  group: "dev"
}
```

Reference to another document



Similar to the object-oriented world
Denormalization!

Nesting, arrays

- Normalization, instead of join

- 1-1 connection

- > Nesting

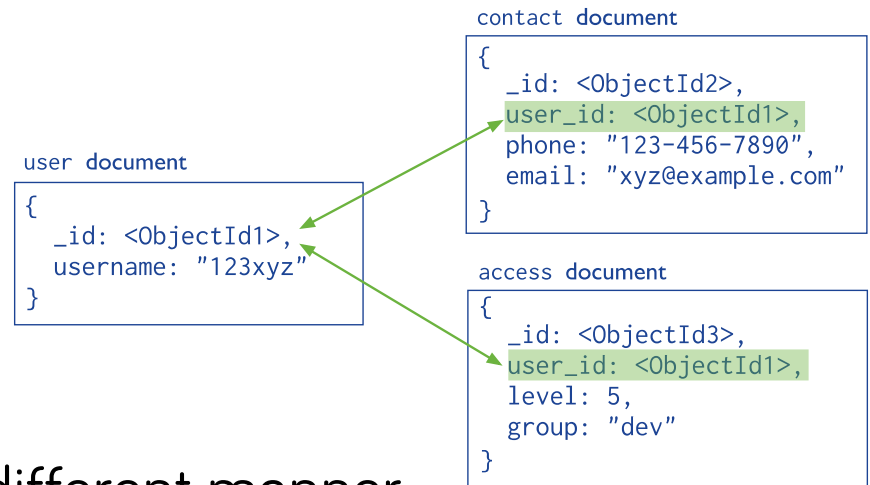


- 1-many connection

- > Embedding as array

- > Or reference

- There is no join, but in a different manner



Transactions

- Usually not supported in this world
- Supported since MongoDB 4.0 (not so long ago)

~~• ACID~~

- Atomicity
- Instead of isolation: read/write concern
 - > Not discussed more in this course.
 - > <https://docs.mongodb.com/manual/reference/read-concern/>
 - > <https://docs.mongodb.com/manual/reference/write-concern/>

MongoDB “schema” design

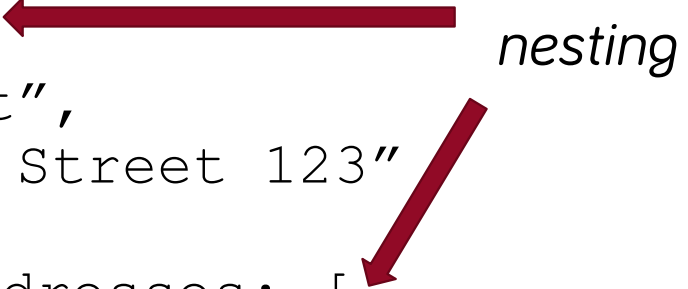
MongoDB “schema” design

- Collections usually yield themselves
- Prefer nesting to alternatives
 - > But consider the document size limit!
- Transactions only when there is no other way

MongoDB schema design - example

- Customers

```
{
  _id: ...,
  name: "John Doe",
  email: john@doe.com,
  mainPlaceOfBusiness: {
    zipCode: 1234,
    city: "Budapest",
    street: "Place Street 123"
  },
  additionalShippingAddresses: [
    { zip:..., city:..., street:... },
    ...
  ]
}
```



nesting

MongoDB schema design - example

- Order


```
{
  _id: ...,
  date: ...,
  customer: ObjectId(...) ← reference
  items: [
    { productId: 111, price: 222, piece: 3 },
    { productId: 111, price: 222, piece: 3 },
    ...
  ],
  shippingAddress: { zip: ..., city: ... }
}
```

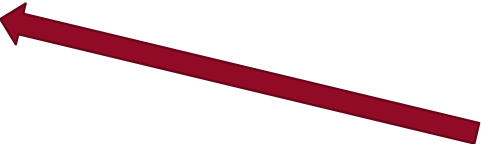
← denormalization

MongoDB schema design - example

- Product

```
{
  _id: ...,
  name: Sofa,
  price: 1234,
  vatPercentage: 27,
  shippingInformation: {
    packagePieces: 2,
    packageSize: [
      {width: 100, height: 45},
      {width: 250, height: 35}
    ]
  }
}
```

 *denormalization*

 *nesting*

Transactions: SQL - NoSQL

- Typical SQL
 - > ACID: Atomicity, Consistency, Isolation, Durability
 - > If it were a single node
- Typical NoSQL
 - > BASE: Basically Available, Soft State, Eventual Consistency
 - > Giving up consistency for availability
 - > Scenario: e.g. number of Facebook likes, name change, etc.

Transactions in MongoDB

- Updating a document is atomic
- It supports distributed transactions
 - > Even with ACID properties
 - > Performance degrades significantly!
- Isolation levels "instead" of concerns :
 - > <https://docs.mongodb.com/manual/reference/read-concern/>
 - > <https://docs.mongodb.com/manual/reference/write-concern/>

SQL vs NoSQL

- SQL challenges
 - Size / rigid data model / **scaling** / **availability**
- DB-level solutions introduced by NoSQL can present new challenges in the other layers
 - Denormalization -> consistency maintenance
 - Scaling -> transactional constraints
 - Semi-structured data -> Type information is lost
 - How is data displayed on the user interface?
 - What data does the API provide to additional layers and external services?
 - How do we match the code to several scheme versions?

MongoDB CRUD operations in .NET

MongoDB protocol

- „Wire protocol”: TCP/IP based binary
- Client applications, drivers for popular platforms
 - > <https://docs.mongodb.com/ecosystem/drivers/>
- Both request and response are JSON documents
 - > It should be thought of like the SQL language, only it is platform specific
 - > Certain drivers provide this level
 - > Driver compatible with more advanced languages (e.g. C#).

MongoDB "raw" language

```
db.inventory.find( {  
    $and: [  
        {  
            price: { $ne: 1.99 }  
        },  
        {  
            price: { $exists: true}  
        }  
    ]  
})
```

- For simplicity, we will NOT show the "raw" JSON queries, only the .NET API

MongoDB .NET

- MongoDB.Driver
- <https://www.nuget.org/packages/mongodb.driver>

→ Check the lecture notes



Code – database mapping

- Code generation
 - > Based on schema – if exists
 - > Based on sample jsons
- Manually – code first

```
public class Product
{
    public ObjectId Id { get; set; } // ez lesz az elsődleges kulcs helyett az _id
    public string Name { get; set; }
    public float Price { get; set; }
    public int Stock { get; set; }
    public string[] Categories { get; set; } // tömb értékű mező
    public VAT VAT { get; set; } // beágyazást alkalmazunk
}

public class VAT // mivel ez beágyazott entitás, így nem adunk neki egyedi azonosítót
{
    public string VATCategoryName { get; set; }
    public float Percentage { get; set; }
}
```

Customization

- Pascalcasing is converted by the built-in convention
 - > We can also create our own conventions
- Customization with attributes

```
public class Product
{
    // _id mezőre képződik le
    [BsonId]
    public ObjectId Azonosito { get; set; }

    // megadhatjuk a MongoDB dokumentumban használatos nevet
    [BsonElement("price")]
    public string Ar { get; set; }

    // kihagyhatunk egyes mezőket
    [BsonIgnore]
    public string NemMentett { get; set; }
}
```

Queries

- We perform operations on collections

```
var collection = db.GetCollection<Product>("products");
```

- Find

> A lambda expression

```
collection.Find(x => x.Price < 123 && x.Name.Contains("red"));
```

> Raw MongoDB query in its own json language

```
collection.Find(  
    Builders<Product>.Filter.And(  
        Builders<Product>.Filter.Lt(x => x.Price, 123),  
        Builders<Product>.Filter.Regex(x => x.Name, "/red/s"  
    )  
);
```

```
{  
  "price": {  
    "$lt": 123.0  
  },  
  "name": "/red/s"  
}
```

Operators

- Filters only apply to constant values

```
collection.Find(x => x.Price != 123);  
collection.Find(Builders<Product>.Filter.Ne(x => x.Price, 123));
```

- Null / non-existent key

```
collection.Find(x => x.VAT != null);  
collection.Find(Builders<Product>.Filter.Exists(x => x.VAT));
```

- Array field filtering

```
// azon termékeket, amelyek a jelzett kategóriában vannak  
collection.Find(Builders<Product>.Filter.AnyEq(x => x.Categories, "Labdák"));
```

- Sorting, pagination

```
collection.Find(...)  
    .Sort(Builders<Product>.Sort.Ascending(x => x.Name))  
    .Skip(100).Limit(100);
```

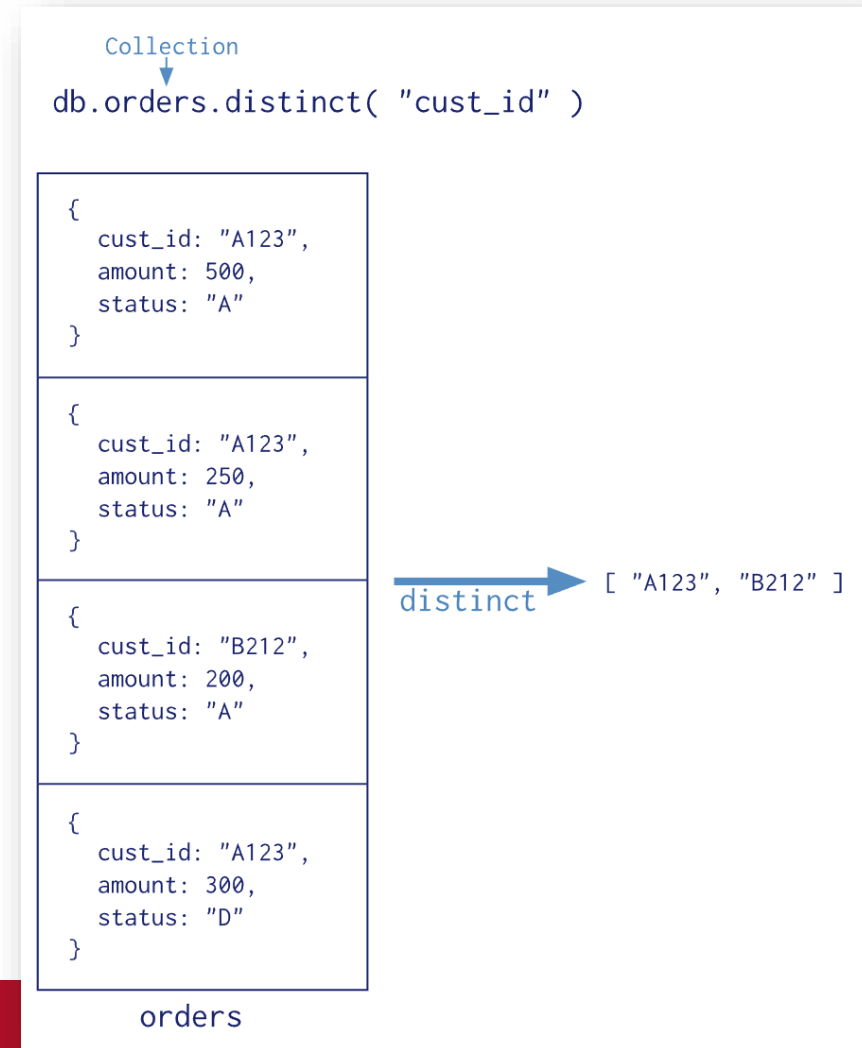
Result processing

- Listing
 - > . ToList()
- Query one element
 - > . First() / FirstOrDefault() / Single() / ...
- Cursor
 - > We process the result continuously

```
var cur = collection.Find(...).ToCursor();  
while (cur.MoveNext()) // kurzor léptetése  
{  
    foreach (var t in cur.Current) // a kurzor aktuális eleme nem egy dokumentum  
    { ... }  
}
```

Aggregation pipeline

- Processing multiple documents on the server side
- Single Purpose
 - > distinct
 - > count



Generic aggregation

- Several actions can be defined one after the other
 - > filtering, grouping, counting, ...

```
// A "Labdák" kategóriába tartozó termékek
// ÁFA kulcs szerint csoportosítva
foreach (var g in collection.Aggregate()
    // szűrés
    .Match(Builders<Product>
        .Filter.AnyEq(x => x.Categories, "Labdák"))
    // csoportosítás
    .Group(x => x.VAT.Percentage, x => x)
    .ToList())
{
    Console.WriteLine($"VAT percentage: {g.Key}");
    foreach (var p in g)
        Console.WriteLine($"  \tProduct: {p.Name}");
}
```

"Join" in mongodb

- Left Outer Join can be requested in an aggregation operation with the \$lookup function
 - > You can even join array elements
- Mongodb is not made for this
- If we were to use too many lookups, think again
 - > Scheme
 - > Is this the right tool for the job?

```
{
  $lookup:
  {
    from: <collection to join>,
    localField: <field from the input documents>,
    foreignField: <field from the documents of the "from" collection>,
    as: <output array field>
  }
}
```


Insert

- Simple object instantiation
- We do not fill in the ID
- **InsertOne / InsertMany**

```
var newProduct = new Product
{
    Name = "Alma",
    Price = 890,
    Categories = new[] { "Gyümölcsök" }
};
collection.InsertOne(newProduct);

Console.WriteLine($"Beszúrt rekord id: {newProduct.Id}");
```

Delete

- Delete based on a filter condition
 - > One or more documents
 - > **DeleteOne / DeleteMany**

```
var deleteResult = collection.DeleteOne(x => x.Id == new ObjectId("..."));  
Console.WriteLine($"Törölve: {deleteResult.DeletedCount} db");
```

Replace entire document

- Replace a document based on filtering
 - > We can even replace the ID
 - > ReplaceOne / FindOneAndReplace

```
var replacementProduct = new Product
{
    Name = "Alma",
    Price = 890,
    Categories = new[] { "Gyümölcsök" }
};
var replaceResult = collection.ReplaceOne(
    x => x.Id == new ObjectId("..."), replacementProduct
);

Console.WriteLine($"Módosítva: {replaceResult.ModifiedCount}");
```

Modify fields

- We modify the specific fields of the retrieved document

> UpdateOne / UpdateMany

```
collection.UpdateOne(  
  filter: x => x.Id == new ObjectId("..."),  
  update: Builders<Product>.Update.Set(x => x.Stock, 5));
```

```
collection.UpdateOne(  
  filter: x => x.Id == new ObjectId("..."),  
  update: Builders<Product>.Update  
    // raktárkészlet legyen 5  
    .Set(x => x.Stock, 5)  
    // mai dátumot beírjuk, mint a frissítés ideje  
    .CurrentDate(x => x.StockUpdated)  
    // töröljük a frissítendő jelzést  
    .Unset(x => x.NeedsUpdate)  
);
```

Insert or modify

- Upsert (update / insert) operation
- We either have it or create it
 - > Modifications can also be made in this way
- Useful for concurrent access

```
collection.ReplaceOne(  
    filter: x => x.Id == new ObjectId("..."),  
    replacement: replacementObject,  
    options: new UpdateOptions() { IsUpsert = true });
```

Field operations

- Set: setting the field value;
- SetOnInsert: like Set, but only executed when a new document is inserted;
- Unset: delete a field (remove the key and value from the document);
- CurrentDate: enter the current date;
- Inc: increase value;
- Min, Max: replacement of the field value, if the specified value is lower/higher than the current value of the field;
- Mul: multiplication of value;
- PopFirst, PopLast: remove first/last element from array;
- Pull: removing a value from an array;
- Push: adding a value to an array at the end (more options in the same operator: sorting an array, keeping the first n elements of an array);
- AddToSet: add a value to an array if it does not already exist.