

Java Persistence API

Imre Gábor

Q.B224

gabor@aut.bme.hu



Automatizálási és
Alkalmazott
Informatikai Tanszék

Outline

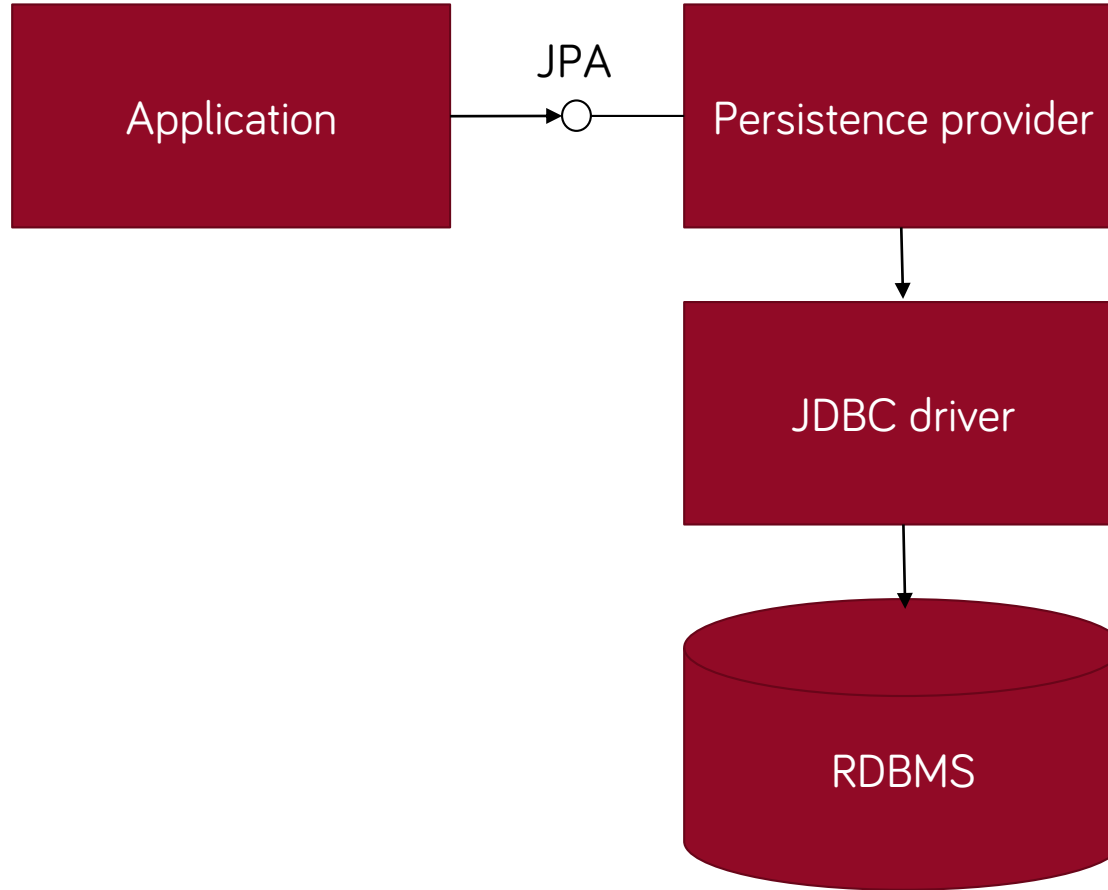
- Overview
- O-R mapping with annotations
- The persistence context
 - > Lifecycle of entities
 - > Database synchronization
 - > Queries
- Criteria API
- Inheritance
- Relationships

Overview

General features


- Standard ORM API in the Java world
- Current version: **JPA 3.2**
- Specifies only interfaces → several implementations (persistence provider – P.P.)
 - > E.g. *Hibernate*, *EclipseLink*, *OpenJPA*
- JPA entity: a class whose instances are stored in a relational database by JPA
- Defined in package **javax.persistence**
 - > Renamed to `jakarta.persistence` after 2.2

JPA architecture



Java Enterprise Edition

- JPA can be used in Java SE as well, by adding the jar files of the persistence provider to the application
- Since Java EE5 each Java EE application server contains a JPA persistence provider → no need to add JPA jar files to Java EE applications
 - > Java EE: Java Enterprise Edition, adds support for server side development of enterprise-scale applications
 - > N-tier architecture
 - > Further Java EE technologies (EJB, JTA) ease the use of JPA



O-R mapping with annotations

O-R mapping with annotations

- Mandatory properties of JPA entities
 - > No-arg constructor
 - > **@Entity** annotation
 - > Primary key attribute annotated with **@Id**
 - The primary key value can be generated via **@GeneratedValue**
 - Composed primary keys are supported via **@IdClass** or **@EmbeddedId**
- Persistent attributes are exposed as getter/setter methods

Entity example

@Entity

```
public class Employee {
```

@Id

```
    private Integer id;
```

```
    private String name;
```

```
    private Date birthDate;
```

```
    // ... getters, setters
```

```
}
```

Customizing the O-R mapping

- Specifying the details of O-R mapping (table, column names) is optional: default values are the names of the class/attributes
- On the class: **@Table(name="MyTable")**
- On the attributes: **@Column(name="MyColumn")**
- XML configuration file can substitute annotations, but rarely used

Entity attribute types

- Primitive types and their wrappers
- String, char[], Character[]
- BigInteger, BigDecimal
- java.util.Date, java.util.Calendar, java.sql.Date, java.sql.Time, java.sql.Timestamp
- byte[], Byte[]
 - > **@Lob**
- Enum
- Other entities or collections of entities, or even non-entities
 - > See Relationships
- Embedded classes
- Prevent an attribute from being stored in DB: **@Transient**

Embedded class

- A class that encapsulates some persistent attributes, but is not an entity on its own

@Embeddable

```
public class EmploymentPeriod {  
    Date startDate;  
    Date endDate;  
    // ... getters, setters  
}
```

- Using an embedded class inside an entity

```
@Entity public class Employee {  
    //...  
    @Embedded  
    @AttributeOverrides({  
        @AttributeOverride(name="startDate", column=@Column("EMP_START")),  
        @AttributeOverride(name="endDate", column=@Column("EMP_END"))  
    })  
    private EmploymentPeriod empPeriod;  
}
```

Persistence unit

- In, JPA, the persistence provider handles persistence units (P.U.)
- P.U.: a set of entities stored in the same database
- Persistence unit(s) are defined in the persistence.xml config file, packaged under the META-INF folder of the jar file containing the entity classes

```
<persistence>
```

```
  <persistence-unit name="MyPU">
```

```
    <jta-data-source>jdbc/MyDB</jta-data-source>
```

```
    <class>com.xy.Employee</class>
```

```
    <class>com.xy.Company</class>
```

```
  </persistence-unit>
```

```
</persistence>
```

Name of the P.U. Multiple persistence-unit tags are needed, when accessing multiple databases

JNDI name of the database

The entity classes. Not necessary to list them in case of Java EE.

JNDI

- Java Naming and Directory Interface
- Java API for accessing naming and directory services
- A name service can store an object registered with a name, by which the object can be looked up later
- Java EE application servers always contain a JNDI provider, for these typical use cases:
 - > Components can look up each other by name
 - > External resources (e.g. DB, message queue, SMTP server) can be looked up by name

Accessing databases in Java EE environment

- Drawbacks of using **DriverManager.getConnection()**:
 - > Developers have to know the details of the connection (DB vendor, server, port, name of DB)
 - > The open connection cannot be reused efficiently
- Solution: **DataSource** interface
 - > Operation (e.g. server administrator) registers the details of the database with a JNDI name in the application server's name service
 - > Developer looks up the datasource via JNDI name
 - > Implements connection pooling, i.e. closing a connection requested from a datasource does not closes it physically, just marks it as free (reusable at other connection requests)

Typical JDBC code in Java EE environment

@Stateless

```
public class MyBean {
```

```
    @Resource(lookup="mydb")
```

```
    DataSource ds;
```

```
    public void myMethod() {
```

```
        try(Connection conn = ds.getConnection()) {
```

```
            ...
```

```
        } catch(Exception e){...}
```

```
    }
```

```
}
```

- Due to **@Resource** the application server performs a JNDI lookup (dependency injection)
- If we would need to do it programmatically:

```
DataSource ds = (DataSource) new InitialContext().lookup("mydb");
```

- JDBC driver class, user/pass, JDBC URL defined in the configuration file of the application server, under **mydb** name

The persistence context

Persistence context (P.C.)

- A set of in-memory entities handled by the persistence provider
- It is the connection between the in-memory entities and the database, we manipulate our entities through the P.C.
- A P.C. is accessed through the EntityManager interface, e.g.

```
EntityManagerFactory emf =  
Persistence.createEntityManagerFactory("MyPU");  
EntityManager em = emf.createEntityManager();  
em.getTransaction().begin();  
em.persist(new Employee(12345, "Gabor"));  
em.getTransaction().commit();  
em.close();  
emf.close();
```

Persistence context

- An EntityManager references one P.C., each entity with a given primary key is contained at most once in a P.C.
- Managing the lifecycle of an EntityManager is challenging:
 - > When should we open it?
 - > When should we close it?
 - When calling another method, it would make sense, if that method could use the same EntityManager reference, so it is needless to close it → it could be a member variable
 - But if it is a member variable, and we call a method in another class, it would make sense if the same EntityManager could be seen there → should it be a method argument? In each method that wants to manipulate entities?
- Solution: dependency injection with Java EE (EJB) or Spring

Managed persistence context

@Stateless

```
class PersonService {  
    @PersistenceContext  
    EntityManager em;  
  
    public void createEmployee(){  
        em.persist(new Employee(12345, "Gabor"));  
    }  
}
```

- Condition: PersonService cannot be directly instantiated with **new**, we should ask a container to give us an instance of PersonService, and that container performs the injection of the proper EntityManager instance
- Both EJB (@Stateless) and Spring (@Service) are capable of handling such *managed persistence contexts*

Managed persistence context

- The EntityManagerFactory
 - > Only one instance is created, when starting the application
 - > We can inject it via **@PersistenceUnit**, but it is seldom needed
- The EntityManager member variable, annotated with **@PersistenceContext**
 - > is created at the beginning of the transaction
 - > is closed at the end of the transaction
 - > if methods in several classes are called in the same transaction, the different EntityManager member variables will access the same P.C.

EntityManager

- The interface for manipulating entities
- 3 types of methods:
 - > Handle the lifecycle of entities
 - > Synchronize entities with DB
 - > Query entities

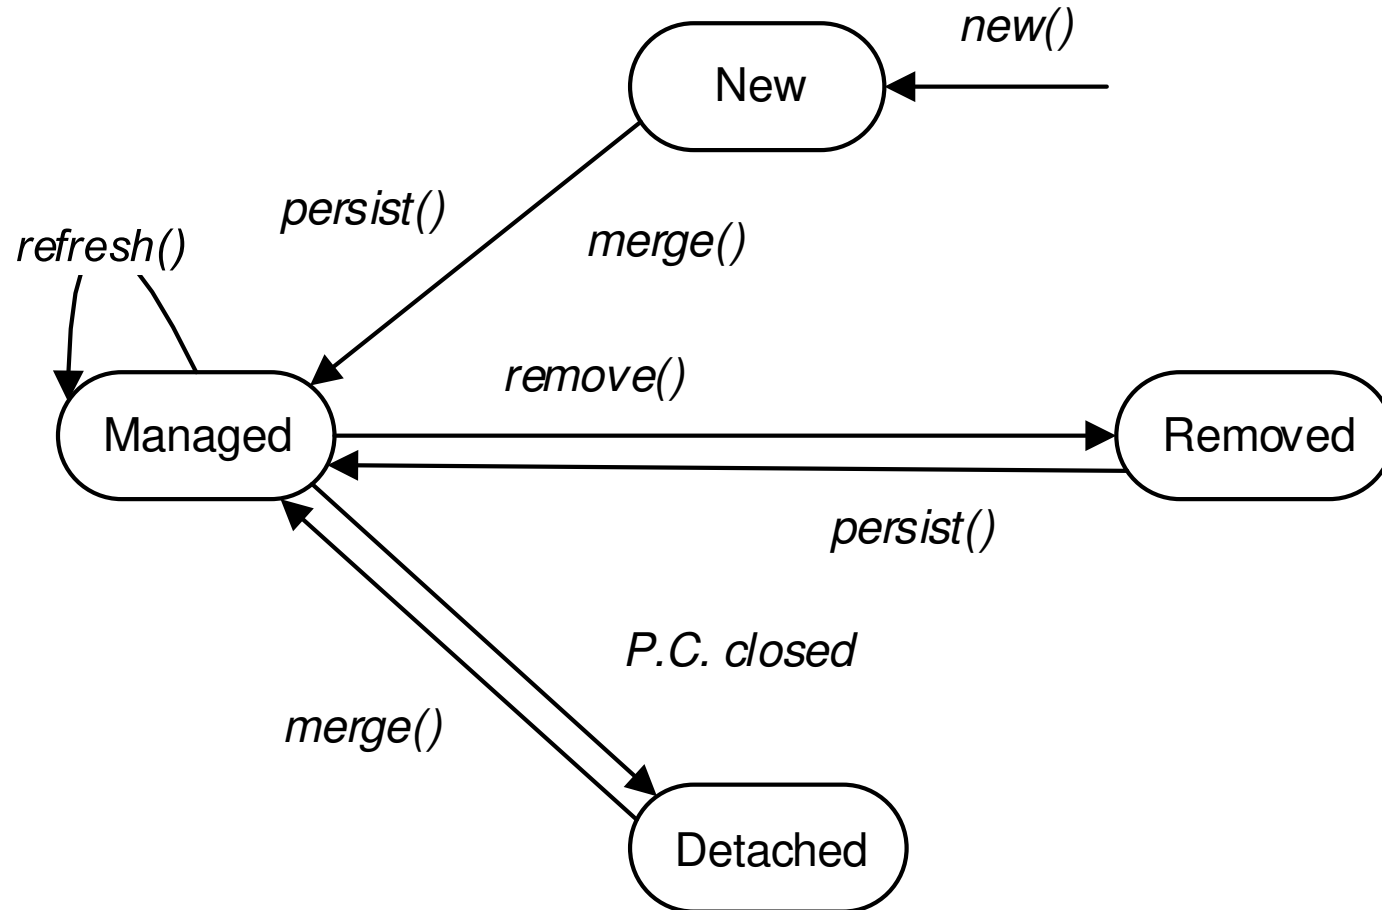
The persistence context

The lifecycle of entities

Entity states

- new: created via new, exists only in memory, is not saved to DB
- managed: exists in the DB, and is contained in a P.C. The in-memory modifications of the entity are saved into DB when the flush() method is called on the EntityManager. (This method is called automatically at transaction commit)
- detached: exists in the DB, but is not contained in the P.C.
- removed: contained in the P.C., but it is marked for removal, will be deleted from DB at the end of the transaction

The lifecycle of entities



The lifecycle of entities

- Making a new entity managed:
 - > **persist()**: tries to INSERT, throws an Exception for existing primary key
 - > **merge()**: performs an SQL UPDATE for existing primary key (and INSERTS for non-existing PK)
- The *return value* of **merge()** is the managed instance
- An entity becomes detached
 - > After clearing the P.C. via **em.clear()**;
 - > After closing the P.C. via **em.close()**;
 - > After detaching it individually via **em.detach(entity)**;

The persistence context

Database synchronization

Database synchronization

- The persistence provider uses 2 methods of the **EntityManager** to synchronize entity data with the database:
 - > **flush()**: writes all the modifications of the managed entities of the P.C. into the DB
 - Rarely called explicitly, because transaction commit calls it automatically
 - It is also called by default before running a query
 - > **refresh(entity)**: reads the most up to date state of the entity from the DB

The persistence context

Queries

Queries

- Multiple possibilities, all through EntityManager
- Query for an entity by its primary key:
 - > `<T> T find(Class<T> entityClass, Object primaryKey)`
- Custom query, dynamically created:
 - > Use JPQL : `public Query createQuery(String jpqlString)`
 - JPQL (Java Persistence Query Language) is a language similar to SQL, but returns complete entity instances
 - E.g. `"SELECT e from Employee e WHERE e.name = :name"`
 - > Use native SQL: `public Query createNativeQuery(String sqlString)`
- Custom query, defined statically, identified by name
 - > `public Query createNamedQuery(String nameOfQuery)`
 - > The JPQL text of the query is defined in a `@NamedQueries` annotation on an entity

`@NamedQueries({`

`@NamedQuery(name="Employee.findAll",
 query="SELECT e FROM Employee e")`

`})`

`@Entity`

`public class Employee { ...}`

Queries

- What can we do with a **Query**?
 - > **setParameter**: by name or by index
 - > Support for paging: **setMaxResult**, **setFirstResult**
 - > Actually running the query:
 - **getSingleResult** (Exception for 0 or more than 1 result)
 - **getResultList** (empty list for 0 results)
 - **executeUpdate** (in case of UPDATE, DELETE)
- Some important JPQL features
 - > Bulk delete, update
 - > JOIN, GROUP BY, HAVING, subquery
 - > Parameters: numbered (?1, ?2, ..) or named (:mypar)
 - > Projection
 - returning only some attributes as List<Object[]>
 - or return custom new objects

JPQL examples

- **SELECT o FROM Order o WHERE o.shippingAddress.state = 'CA'**
 - > Selects the Order entites that have to be shipped to California
- **SELECT DISTINCT o FROM Order o JOIN o.lineItems l WHERE l.shipped = FALSE**
 - > Selects the Order entities that contain not yet shipped orderItems

Criteria API

Criteria API

- A typesafe, object-oriented alternative of the string based JPQL for creating queries (since JPA 2.0), e.g.

```
CriteriaBuilder cb = em.getCriteriaBuilder();
```

```
CriteriaQuery<Employee> cq =  
    cb.createQuery(Employee.class);
```

```
Root<Employee> emp = cq.from(Employee.class);
```

```
cq.select(emp);
```

```
cq.where(cb.equal(emp.get("lastName"), "Smith"));
```

```
TypedQuery<Employee> query = em.createQuery(cq);
```

```
List <Employee> rows = query.getResultList();
```

Criteria API

- More code than in case of JPQL, but the compiler protects us against type errors and typos (this version not yet completely)
- This last solution contains a string based attribute ("lastName"), this can be eliminated by using the Metamodel API
- The metamodel of the persistence unit contains metadata about the entities (and embedded classes)
- Can be generated via tools of the persistence provider, but could be written by hand as well

Criteria API

@Entity

```
public class Employee {  
    @Id Long id;  
    String firstName;  
    String lastName;  
    Department dept;  
}
```

@StaticMetamodel(Employee.class)

```
public class Employee_ {  
    public static volatile SingularAttribute<Employee, Long> id;  
    public static volatile SingularAttribute<Employee, String> firstName;  
    public static volatile SingularAttribute<Employee, String> lastName;  
    public static volatile SingularAttribute<Employee, Department> dept;  
}
```

Criteria API

- Our criteria queries can be completely type-safe with the use of the metamodel class, e.g.:

```
CriteriaQuery<Employee> cq =  
    cb.createQuery(Employee.class);  
  
Root<Employee> emp = cq.from(Employee.class);  
cq.select(emp);  
cq.where(cb.equal(  
    emp.get(Employee_.lastName), "Smith"));  
  
TypedQuery<Employee> query = em.createQuery(cq);  
List<Employee> rows = query.getResultList();
```

Inheritance

Mapping inheritance in JPA

- Implementing the inheritance hierarchy of classes
 - > **extends** keyword
 - > **@Entity** on each entity
 - > **@Id** attribute only on the topmost superclass
- The mapping of the inheritance is defined at the topmost superclass, via annotation:
 - > **@Inheritance(strategy = SINGLE_TABLE)** → Mapping into a single table
 - > **@Inheritance(strategy = JOINED)** → Mapping all classes
 - > **@Inheritance(strategy = TABLE_PER_CLASS)** → Mapping real classes
 - Not mandatory for a P.P. To support it, because supporting polymorphism is cumbersome

Discriminator column

- The table itself determines the Java type only in case of **TABLE_PER_CLASS**
- In case of **SINGLE_TABLE** and **JOINED** an additional column is needed to store the type → discriminator column
- The name of this column is by default **DTYPE**
 - > But can be overridden in the topmost superclass:
@DiscriminatorColumn(name="mycolumn")
- The value stored in this column is by default the name of the entity
 - > But can be overridden at each entity:
@DiscriminatorValue("mytype")

Other means of inheritance

- An entity can be a subclass of a non-entity class
 - > Question: should the attributes in the non-entity parent class be persistent in the entity subclasses?
 - By default not
 - But marking the non-entity superclass with **@MappedSuperClass**
→ yes
- Non-entity classes can inherit from entity classes
 - > It won't have a corresponding table, cannot be used in queries or passed to EntityManager methods
- An entity class can be abstract
 - > Cannot be instantiated, but can be mapped to a table (in case of JOINED inheritance), can be used in queries

Relationships

Mapping relationships

- The type of the member variable representing the relationship
 - > Either an entity
 - > Or **Collection**, **Set**, **List** or **Map** of entities
- The member variable must be annotated based on the multiplicity:
 - > **@OneToOne** (Implemented with a foreign key in DB)
 - > **@OneToMany** (Implemented with a foreign key in DB)
 - > **@ManyToOne** (Implemented with a foreign key in DB)
 - > **@ManyToMany** (Implemented with a join table in DB)
- Further possible annotations:
 - > **@JoinColumn** (when omitting, default column name is used)
 - > **@JoinTable** (in case of many-to-many, when omitting, default table name is used)
 - > **@OrderBy** (in case of a List relationship)
 - > **@MapKey** (in case of a Map relationship)

Mapping relationships

- Based on direction:
 - > Uni-directional
 - > Bi-directional
 - The two sides are kept consistent by the developer!!!
 - The two sides of the relationship have to be mapped to each other, e.g.

- On the owner side (Employee):

@ManyToOne

@JoinColumn(name="company_id")

private Company company;

- On the other side (Company):

@OneToMany(mappedBy="company")

private Collection<Employee> employees;

- The relationship has always one owner side. The other side references to it via mappedBy

Collection typed attributes with non-entity elements

- Available since JPA 2.0, for base types or embedded classes
- Separate table for the collection elements, with foreign key pointing on the parent entity
- But the table of the collection has neither primary key, nor corresponding entity
- Can be customized via **@ElementCollection** és **@CollectionTable**

@ElementCollection example

```
public enum FeatureType { AC, CRUISE, PWR, BLUETOOTH, TV, ... }
```

@Embeddable

```
public class ServiceVisit {  
    @Temporal(DATE)  
    Date serviceDate;  
    String workDesc;  
    int cost;  
}
```

@Entity

```
public class Vehicle {  
    ...
```

@ElementCollection example

@Entity

```
public class Vehicle {
```

```
    @Id int vin;
```

```
    @ElementCollection
```

```
    @CollectionTable(name="VEH_OPTNS")
```

```
    Set<FeatureType> optionalFeatures;
```

```
    @ElementCollection
```

```
    @CollectionTable(name="VEH_SVC")
```

```
    @OrderBy("serviceDate")
```

```
    List<ServiceVisit> serviceHistory;
```

```
    ...
```

```
}
```

Relationships

Fine-tuning

Cascade

- All 4 relationship annotations accept a cascade parameter, e.g.

```
@OneToMany(cascade={  
    CascadeType.PERSIST, CascadeType.MERGE  
})
```

- Possible values: PERSIST, MERGE, REMOVE, REFRESH, ALL
- It defines the EntityManager operations that should be automatically called on the related entity instances
- Default: no cascade

Fetch

- All 4 relationship annotations accept a fetch parameter, e.g.

@OneToMany(fetch=FetchType.LAZY)

- It defines, whether the related entities should be automatically loaded when loading the entity itself
- LAZY: loaded only when actually accessing the relationship → no needless memory usage, but +1 query
 - > P.P. can ignore it!
- EAGER (default, except for OneToMany és ManyToMany): loaded immediately, together with the referring entity → faster, but may consume memory in vain
 - > Mandatory for the P.P., not ignorable
- Possible fine-tuning:
 - > Eager fetch defined in the query:
SELECT c FROM Customer c LEFT JOIN FETCH c.orders
- Problem related to fetch
 - > Entity becoming detached without the lazy-fetched relationships loaded in memory → trying to access them in detached state throws Exception